

REALITY UNDER INQUIRY

Toward an Epistemology and Ontology of Science

Scott Brodie Forsyth

REALITY UNDER INQUIRY

Toward an Epistemology and Ontology of Science

Scott Brodie Forsyth

This is an original philosophical work. Historical and scientific claims are distinguished from interpretation, argument, and speculation. References are given without invented quotations or page numbers; edition-specific pagination should be checked against the copy used by a reader.

Copyright © 2026 Scott Brodie Forsyth. All rights reserved.

Abstract

Science aims at knowledge of a reality that does not depend on our wishes, languages, or institutions. Yet scientific knowledge is produced by historically situated people through instruments, models, measurements, statistical methods, and social systems that can both expose and conceal the world. This book asks whether objectivity is possible under those conditions, and if so what it means.

The answer developed here is *Reflexive Constraint Realism* (RCR). Its ontological thesis is that a mind-independent world contains entities, processes, and relations that constrain what can happen. Its epistemological thesis is that claims become objective to the extent that they survive independent routes of constraint: reliable measurement, controlled variation, intervention, cross-context transport, model comparison, critical disagreement, and organized attempts to find error. The account is realist without being naive, social without being relativist, and fallibilist without treating every belief as equally good.

RCR distinguishes four layers of scientific success. Correspondence concerns whether a claim is about what exists or occurs. Structural adequacy concerns whether a representation preserves relevant relations. Interventional adequacy concerns whether manipulating a variable changes outcomes as the model says. Pragmatic adequacy concerns whether a representation supports reliable inquiry and action. These layers can come apart. A model may be useful while idealized; an entity may be real while a theory about its deeper nature is wrong; and a prediction may succeed for the wrong causal reason. The book therefore proposes graded ontological commitment, explicit scope conditions, and a social norm of public corrigibility.

The framework is reconstructed through logical empiricism, falsificationism, Kuhn, Lakatos, Feyerabend, Bayesian and error-statistical approaches, inference to the best explanation, realism and anti-realism, pragmatism, social and feminist epistemology, naturalized epistemology, causal inference, model-based science, and debates about explanation, laws, reduction, emergence, values, replication, uncertainty, and demarcation. It is then tested against cases in physics, biology, medicine, psychology, economics, climate science, and the social sciences. The final chapters state the the-

ory's axioms, inferential rules, limits, objections, and an empirical research programme.

Preface

The philosophy of science often appears as a sequence of positions: empiricism, falsificationism, paradigms, research programmes, Bayesianism, realism, anti-realism, and so on. That sequence is useful, but it can give the false impression that the central task is to choose a camp. The more demanding task is to understand why each camp saw something important, where it overreached, and what a coherent account of objectivity would have to preserve.

This book is written in that spirit. It does not treat science as a single algorithm. Nor does it treat the diversity of scientific practices as evidence that there are no standards beyond local agreement. A laboratory experiment, a climate model, a historical dataset, a randomized medical trial, and a causal graph do not establish claims in the same way. They can nevertheless be judged by shared questions: What is being claimed? What would count against it? Which parts of the result depend on instruments, models, or social assumptions? Can the claim travel to a new setting? What would an independent critic need in order to find an error?

The framework proposed here is original in its organization and argumentative role, not in the sense that every component has no precedent. The literature already contains powerful ideas about intervention, structure, mechanisms, social criticism, error, and realism. The contribution is to place these ideas inside one account of objective inquiry, to specify their relations, and to make their limits explicit.

The reader will find some formal notation. It is used as a tool for precision, not as a display of technical status. Every symbol is introduced before it is used. Where a historical claim is compressed, the bibliography points to a primary source or a respected scholarly treatment. Where the book advances an interpretation, it says so. Where it speculates, it marks the speculation.

Introduction

Science is at once our most powerful way of learning about the world and a human practice conducted through imperfect concepts, instruments, models, institutions, and values. The central problem of this book is how those two facts belong together. If science is objective, objectivity cannot mean infallibility or a view from nowhere. If science is situated and revisable, revisability cannot mean that evidence is powerless or that every description is equally good.

The argument proceeds in three movements. The first reconstructs the principal traditions in philosophy of science, from empiricism and logical positivism through falsificationism, historical theories of change, Bayesian and explanatory approaches, realism and anti-realism, pragmatism, and social and feminist epistemologies. The second develops an account of observation, models, measurement, causality, explanation, mechanisms, laws, emergence, uncertainty, replication, and values. The third states and tests *Reflexive Constraint Realism* against cases in physics, biology, medicine, psychology, climate science, economics, and the social sciences.

The result is a layered epistemology and ontology. It treats scientific claims as scope-bound achievements whose warrant depends on evidence, invariance, intervention, model transparency, and organized criticism. It also treats their ontological force as graded: a successful pattern does not automatically establish a hidden entity, but convergent structural and causal success can justify increasingly strong commitments. The conclusion is a form of realism disciplined by reflexivity, explicit uncertainty, and the possibility of correction.

Contents

Abstract	iii
Preface	v
Introduction	vii
1 The Problem of Objectivity	1
1.1 Four questions hidden inside “objectivity”	2
1.2 The temptation of a single method	2
1.3 Mind-independence is not observer-independence	3
1.4 The architecture of the argument	4
1.5 What would count as an answer?	5
2 The Empirical and Logical Traditions	7
2.1 Meaning, evidence, and the rejection of pseudo-problems	7
2.2 The appeal of protocol statements	8
2.3 Confirmation and the problem of induction	9
2.4 Theoretical terms are not optional decorations	10
2.5 Holism and the Duhem–Quine problem	10
2.6 What survives	11
3 Falsification, Revolutions, and Research Programmes	13
3.1 Popper’s strongest falsificationism	13
3.2 Kuhn and the structure of scientific change	14
3.3 Lakatos and the time-scale of criticism	15
3.4 Feyerabend and the criticism of methodological monism .	15
3.5 Are revolutions rational?	16
3.6 The lesson for objectivity	17
4 Realism, Anti-Realism, and the Problem of Success	19
4.1 Constructive empiricism	20
4.2 Scientific realism and the no-miracles intuition	20
4.3 The pessimistic meta-induction	21
4.4 Structural realism	21
4.5 Entity realism and intervention	22

4.6	Pragmatism without relativism	22
4.7	A layered account of truth	23
4.8	What realism should commit us to	23
5	Observation, Measurement, and the Making of Data	27
5.1	Theory-ladenness without observational chaos	27
5.2	The measurement path	28
5.3	Operationalism and construct validity	29
5.4	From data to phenomena	29
5.5	Error, uncertainty, and precision	30
5.6	Instrumental realism	30
5.7	Observation in the social sciences	31
6	Models, Explanation, Prediction, and Laws	33
6.1	Models as mediators	33
6.2	Idealization and abstraction	34
6.3	Prediction and explanation	35
6.4	Laws of nature and local regularities	35
6.5	Reduction and emergence	36
6.6	Simulation and computational models	37
6.7	The model's honesty condition	37
7	Causality, Mechanisms, and Complexity	39
7.1	Potential outcomes and the missing counterfactual	39
7.2	Graphs and confounding	40
7.3	Bayesian reasoning about causal claims	41
7.4	Interventionist explanation	41
7.5	Mechanisms and causal structure	42
7.6	Complexity, feedback, and nonlinearity	42
7.7	What causal knowledge does not require	43
8	Social Epistemology, Standpoints, and Values	45
8.1	The social organization of knowledge	45
8.2	Merton's norms and their limits	46
8.3	Longino and transformative criticism	46
8.4	Standpoint theory and situated knowledge	47
8.5	Values in science	48
8.6	When institutions corrupt inquiry	49
8.7	Objectivity as social self-correction	49
9	Uncertainty, Replication, and Demarcation	51
9.1	Probability as rational change	51

9.2	Error-statistical learning	52
9.3	Replication is not a magic word	52
9.4	Uncertainty and the replication ladder	53
9.5	The demarcation problem	53
9.6	Scientism and the misuse of science	54
9.7	Reliability as a system property	54
10	The Ontology of Reflexive Constraint Realism	57
10.1	Generative reality	57
10.2	Eight axioms	58
10.3	What scientific entities are	59
10.4	Structures and relations	60
10.5	Laws as modal constraints	60
10.6	The ontology of social facts	60
10.7	Graded ontological commitment	61
10.8	A layered map of reality	61
11	The Epistemology of Reflexive Constraint Realism	65
11.1	Claims, targets, and alternatives	65
11.2	Evidence as constraint	66
11.3	The inferential rules	66
11.4	Belief revision	67
11.5	Theory comparison	68
11.6	Explanation and evidence are distinct	69
11.7	Objectivity criteria	69
11.8	What makes the epistemology reflexive	70
12	Scientific Progress and Theory Choice	73
12.1	Against simple accumulation	73
12.2	The problem of underdetermination	74
12.3	Convergence and robustness	74
12.4	Simplicity and complexity	75
12.5	The role of explanation	75
12.6	Theory change and retained constraints	75
12.7	Research programmes and the future	76
13	Cases: Testing Reflexive Constraint Realism	79
13.1	Physics: atoms, electrons, and theory change	79
13.2	Biology: the structure of DNA	80
13.3	Medicine: the MRC streptomycin trial	81
13.4	Psychology: replication and context	81
13.5	Climate science: convergence across models and evidence	82

13.6	Economics: policy models and the Lucas critique	83
13.7	Social science: field experiments and discrimination . . .	83
13.8	Cross-case comparison	84
14	Objections, Revisions, and Limits	87
14.1	Objection 1: the framework merely renames realism . . .	87
14.2	Objection 2: constraints are too broad to be explanatory	88
14.3	Objection 3: convergence can be correlated failure	88
14.4	Objection 4: intervention is theory-laden	89
14.5	Objection 5: Bayesian updating makes objectivity subjective	89
14.6	Objection 6: social criticism cannot reach mind-independent truth	90
14.7	Objection 7: standpoint theory becomes identity essentialism	90
14.8	Objection 8: values contaminate objectivity	90
14.9	Objection 9: the commitment ladder is arbitrary	91
14.10	Objection 10: the framework is too demanding for real research	91
14.11	Objection 11: complexity defeats stable constraints . . .	92
14.12	Objection 12: the framework risks relativism	92
14.13	Objection 13: the framework risks scientism	92
14.14	The revised framework	93
	Conclusion: A Programme for Objective Inquiry	95
	The systematic statement	95
	What RCR avoids	97
	Unresolved problems	97
	Possible empirical implications	98
	A research programme for further development	99
	Final proposition	99
	Glossary	101
	Editorial and Factual Review	107
	Factual and source review	107
	Argument review	107
	Reference review	109
	Language and editorial review	109
	Visual and typesetting review	109
	Bibliography	111
	Detailed Index	119

The Problem of Objectivity

Chapter guide. This chapter frames the book's central problem. It distinguishes a mind-independent world from the human practices by which claims about that world are produced. It also introduces the four questions that any adequate philosophy of science must answer: what exists, how evidence supports belief, what scientific representations mean, and how institutions affect inquiry.

Science is successful in an ordinary and powerful sense. It predicts eclipses, builds vaccines, identifies molecules, detects gravitational waves, estimates the effects of policies, and changes what human beings can do. Yet success alone does not settle what science knows. A horoscope can sometimes make a vague prediction that appears correct. A simple model can forecast well while misrepresenting the mechanism that generates the forecast. An institution can produce reliable measurements while excluding the people most affected by its research questions. Scientific success is evidence, but it is not a complete theory of evidence.

The question of this book is therefore not whether science works. It is how scientific work can yield objective knowledge when the work is performed by finite, situated, interested, and fallible beings. We use concepts to identify objects. We use instruments that embody theories. We compare observations by statistical models. We inherit categories from institutions and languages. If these conditions mediate every route to the world, why should the outcome be anything more than a socially stabilized description?

The answer cannot be that scientists simply look harder. Nor can it be that scientific claims are objective because a community accepts them. A community can be mistaken, and a dissenting individual can be right. The answer must explain how social practices can improve contact with a world that does not depend on those practices.

1.1 Four questions hidden inside “objectivity”

The word *objectivity* is used for several different achievements. A measurement may be objective because different observers can obtain it. A claim may be objective because it is not tailored to one person’s interests. A method may be objective because it makes its assumptions visible. A theory may be objective because it tracks features of the world. These are related but not identical.

This book separates four questions.

1. **The ontological question:** what, if anything, exists independently of human observation, language, and classification? Is reality made only of observable events, or should science also commit us to atoms, genes, fields, social structures, or mechanisms?
2. **The epistemological question:** how can evidence justify a belief about that reality? What is the role of observation, experiment, probability, causal intervention, explanation, and criticism?
3. **The representational question:** what do theories and models represent? Do they describe entities, mathematical structures, relations, possibilities, instruments, or only the data they organize?
4. **The institutional question:** how do communities, incentives, division of labour, expertise, funding, authority, and exclusion affect the production of knowledge?

A position can answer one question well and another badly. Constructive empiricism gives a disciplined account of belief in observables while remaining agnostic about unobservables. Structural realism identifies an important continuity across theory change but may leave the nature of structure unclear. Social epistemology explains why criticism by others matters but cannot by itself guarantee that a community tracks the world. A complete account must connect the four questions without reducing them to one another.

1.2 The temptation of a single method

The search for a single scientific method has an understandable motivation. If there were one valid procedure, objectivity would be easy to define: apply the procedure and accept its results. But the sciences do not have one uniform procedure. Astronomers infer from light that cannot be manipulated in a laboratory. Geologists reconstruct deep time from traces. Epidemiologists compare populations under ethical constraints.

Particle physicists build detectors whose output depends on sophisticated calibration. Social scientists study agents who respond to being studied. Mathematical scientists prove theorems and use simulations to explore spaces too large for analytic solutions.

This diversity does not imply that science is methodologically anarchic. It implies that standards must be abstract enough to travel across practices without becoming so abstract that they say nothing. The shared standard proposed here is not a recipe but a question:

Constraint question. Which features of the claim are forced by independent interactions with the world, and which features are supplied by the investigator’s choices of concepts, models, instruments, and purposes?

The question can be made precise in different ways. In a randomized trial, independence may come from treatment assignment and outcome measurement. In climate science, it may come from the convergence of observations, physical theory, attribution studies, and model ensembles. In a historical science, it may come from traces that are causally connected to past events and from independent archives. In each case the objectivity of the claim depends on the specific constraints that resist convenient reinterpretation.

1.3 Mind-independence is not observer-independence

A reality can be independent of our concepts while being accessible only through conceptual and technological mediation. A mountain is not created by the word “mountain”, but the boundaries of a geological unit depend on purposes, scale, and theory. The social category of unemployment is not a mere illusion, but its operational definition is an institutional achievement. An electron is not invented by a detector, but the detector’s output is not self-interpreting.

This distinction avoids two bad extremes.

- **Naive realism** treats perception as a transparent window and assumes that scientific categories simply copy natural joints.
- **Relativism** treats the dependence of descriptions on concepts as evidence that there is no fact of the matter beyond local schemes.

The middle position requires a theory of mediation. Concepts, instruments,

and models are not obstacles that must be removed before inquiry begins. They are means of contact. Their quality is tested by what they enable investigators to detect, manipulate, predict, explain, and correct. A microscope does not show “reality without mediation”. It creates a controlled relation in which some features of a specimen become more accessible and other features become invisible. Objectivity is achieved when the mediation is understood well enough that the result does not depend on one unexamined viewpoint.

1.4 The architecture of the argument

The argument proceeds in three stages.

First, the major traditions are reconstructed. Logical empiricism correctly insisted that public evidence and formal clarity matter. Popper correctly emphasized risk and criticism. Kuhn correctly showed that research is organized by historically situated paradigms. Lakatos correctly observed that theories are tested as evolving programmes rather than isolated sentences. Feyerabend correctly warned against turning a successful history into a rigid rulebook. Bayesian epistemology correctly clarifies rational belief revision. Realist and anti-realist traditions each identify real tensions in the relation between successful theories and invisible entities.

Second, the practices of science are examined: measurement, model building, explanation, intervention, mechanisms, social criticism, replication, and uncertainty. These practices reveal that no single relation to truth is sufficient. A good scientific representation may preserve a structure without describing every detail. A causal claim requires more than a fitted curve. A replication can fail because a result was false, because a population changed, or because the original effect was context-sensitive. Institutional criticism can increase objectivity without making evidence a product of power alone.

Third, the book develops *Reflexive Constraint Realism* (RCR), an original framework with an ontology, an epistemology, inferential rules, and explicit limits. The framework holds that the world is mind-independent and generative: it produces regularities, resistances, capacities, and transformations. Science does not mirror the whole world. It constructs representations that isolate and track constraints relevant to a question. A claim is objective when its support is distributed across sufficiently independent routes of inquiry and when the inquiry exposes the conditions under which the claim would fail.




Figure 1.1: The basic direction of RCR: claims are grounded by contact and tested by further constraint. The arrows are not a single linear method.

1.5 What would count as an answer?

An adequate philosophy of science should not merely repeat that science is fallible. It should specify how fallibility and objectivity coexist. It should explain why some revisions are improvements rather than merely changes. It should say when a model's idealization is harmless and when it distorts the target. It should distinguish a useful forecast from a confirmed causal mechanism. It should show how values can guide research without determining empirical conclusions. It should explain why disagreement can be epistemically productive in one setting and corrosive in another.

The framework developed later makes four commitments.

1. Reality is not constituted by inquiry, although inquiry can constitute objects of social practice and categories of measurement.
2. Evidence is relational: its force depends on a claim, an alternative, a measurement process, and a background model. This does not make evidence subjective; it makes its conditions explicit.
3. Scientific representations are selective structures. Their truth is layered rather than all-or-nothing.
4. Objectivity is a property of an inquiry system, not only of an isolated sentence. A system is objective when it makes error discoverable by agents who can challenge one another under transparent, world-sensitive constraints.

These commitments are not asserted as axioms at the beginning because their meaning depends on the history and practice that follow. The next chapters therefore begin with the traditions that made the problem unavoidable.

Chapter summary

The central problem is how a socially and conceptually mediated practice can produce knowledge of a mind-independent world. Objectivity has ontological, epistemological, representational, and institutional dimensions. Science has no single algorithm, but its diverse methods can be evaluated by asking which parts of a claim are constrained by independent interactions with the world. The book's proposed answer, Reflexive Constraint Realism, treats objectivity as stabilized world-tracking under criticism, intervention, and explicit scope conditions.

Objections and discussion.

1. Is a measurement objective if all observers agree but the instrument's calibration is unknown?
2. Can a model be scientifically valuable even if no part of it is literally true?
3. What would count as an independent route of evidence in a historical science where controlled intervention is impossible?

Further reading. For orientation, compare the historical reconstruction in Kuhn (1962), the realist debate in van Fraassen (1980) and Chakravartty (2007), the experimental emphasis of Hacking (1983), and the social account of objectivity in Longino (1990).

The Empirical and Logical Traditions

Chapter guide. Logical empiricism is reconstructed as a family of projects rather than a single doctrine. The chapter then examines its achievements and difficulties: the demand for public evidence, formal clarification, reduction of theoretical claims, probability, confirmation, and the problems of meaning, holism, and theory-ladenness. It closes with the lesson that empiricism supplies a discipline of contact but not a complete account of objectivity.

The label *logical positivism* is often used as if it named a uniform creed. Historically, the movement was more varied. The Vienna Circle, the Berlin group, and associated philosophers disagreed about probability, protocol statements, physicalism, conventionality, and the proper role of metaphysics. The later label *logical empiricism* is often more accurate because several figures abandoned the strongest version of a verification principle and developed sophisticated accounts of theoretical terms, confirmation, and scientific language.

The broad ambition was clear: make scientific claims answerable to experience and make their logical structure explicit. This ambition remains part of any defensible epistemology. The problem is not the desire for evidence. The problem is the idea that evidence can be specified without theoretical, instrumental, historical, or social conditions.

2.1 Meaning, evidence, and the rejection of pseudo-problems

Logical empiricists sought a criterion for distinguishing statements that are empirically meaningful from those that merely appear to make claims. A mathematical statement could be meaningful through formal relations. An empirical statement could be meaningful through its connection to possible

experience. Traditional metaphysical sentences were criticized when they seemed to lack a method of verification or a clear inferential role.

The strongest charitable interpretation is methodological rather than dismissive. The empiricists were not simply saying that everything not immediately observable is nonsense. They were asking what difference a claim makes to possible observation, calculation, explanation, or action. A statement about an electron could be meaningful if it entered laws and experimental procedures that generated testable consequences. A statement about a transcendent entity that made no difference to any possible inquiry would not belong to empirical science.

This motivation can be retained without adopting a strict verification criterion. Scientific meaning is partly a matter of inferential role: what would the claim predict, explain, license, rule out, or change in a measurement procedure? The fact that a term is theoretical does not disqualify it. It requires a map of how the term connects to observation and intervention.

2.2 The appeal of protocol statements

A basic problem in empiricism is the status of the observational basis. Scientific claims are tested against observations, but observations are reported in language and obtained under conditions. Is there a neutral layer of protocol statements that is independent of theory?

Different empiricists proposed different answers. Some emphasized immediate experience; others moved toward physical language, public records, or intersubjective agreement. The movement away from private sense data was a major strength. A scientific observation must be inspectable by others, linked to a procedure, and embedded in a record. A private sensation can be evidence for a person's experience, but it is not yet a public measurement.

Yet public accessibility is not the same as theoretical neutrality. To report that a pointer moved, one needs a concept of a pointer and a rule for reading it. To report a spectral line, one needs an instrument, a calibration procedure, and a model connecting detector counts with wavelength. Observation sentences are not arbitrary, but they are not self-authenticating either.

Empirical contact principle. A scientific observation is not a bare given. It is a publicly inspectable result of a procedure that connects a physical interaction to a claim under stated assumptions.

This principle keeps the empiricist demand for public contact while refusing the fantasy of an uninterpreted foundation. The later chapters will call the stated assumptions a *measurement path*: the chain from target, through instrument and calibration, to recorded datum and interpreted result.

2.3 Confirmation and the problem of induction

Suppose a scientist observes many white swans. The observations support the claim that all swans are white, but they do not entail it. The next swan could be black. This is the familiar problem of induction. The problem is not specific to swans. A theory may fit every observation collected so far and still fail tomorrow.

Logical empiricists responded in several ways. Some treated confirmation as a degree of support rather than proof. Reichenbach developed a probabilistic framework in which inductive rules were justified pragmatically: if a stable frequency existed, a suitable rule would converge toward it. Carnap attempted to formalize degrees of confirmation. Others accepted that scientific inference is ampliative and emphasized the practical success of prediction.

The important lesson is that evidence should change belief without pretending to establish certainty. A rational method needs at least three ingredients:

1. a set of candidate claims;
2. information about how likely the evidence would be under those claims;
3. a rule for representing uncertainty and updating.

Bayesian epistemology provides one such rule. If H is a hypothesis and E is evidence, then Bayes' theorem is

$$\mathbb{P}(H \mid E) = \frac{\mathbb{P}(E \mid H)\mathbb{P}(H)}{\mathbb{P}(E)}.$$

Here $\mathbb{P}(H)$ is the prior probability assigned before E , $\mathbb{P}(E \mid H)$ is the likelihood of observing E if H is true, and $\mathbb{P}(H \mid E)$ is the posterior probability after the evidence. The denominator $\mathbb{P}(E)$ makes the posterior probabilities sum to one across the relevant alternatives.

The equation is a consistency rule, not a machine that discovers the correct prior. If two investigators assign different priors, they may initially disagree

even after observing the same evidence. They may nevertheless converge when the likelihood ratio is sufficiently strong. This observation will matter later: public objectivity cannot be identified with identical starting beliefs, but it can be strengthened by evidence that sharply discriminates among alternatives.

2.4 Theoretical terms are not optional decorations

A strict empiricism sometimes imagines that science could be translated into a language containing only observable predicates. But theoretical terms do more than abbreviate observations. They organize relations among observations, generate new questions, and make interventions possible.

Consider temperature. A thermometer reading is not identical to temperature. The reading is produced by a physical process whose response is calibrated against a scale. The theoretical concept of temperature links different instruments, equilibrium conditions, equations, and predictions. Removing the term would not make the science more empirical; it would make the connections harder to express.

The same point applies to genes, fields, utility, inflation expectations, and neural representations. Some of these terms refer to entities, some to properties, some to model variables, and some to institutional constructs. They should not all receive the same ontological status. But they cannot be evaluated merely by asking whether they are directly observed.

A scientific term earns epistemic standing by the role it plays in a network of constraints. It should connect to measurement, support counterfactuals, survive alternative operationalizations, or make possible reliable intervention. The term becomes suspicious when it is insulated from every possible correction and is used only to redescribe its own successes.

2.5 Holism and the Duhem–Quine problem

An isolated hypothesis rarely entails an observation by itself. A prediction typically depends on auxiliary assumptions about initial conditions, instruments, background laws, numerical approximations, and data processing. When the prediction fails, several components may be responsible. This is the problem associated with Duhem and later Quine.

Quine's critique of the analytic–synthetic distinction extended the point: beliefs form a web, and experience bears on the web as a whole rather than on single statements one at a time (Duhem, 1906; Quine, 1951).

The strongest lesson is epistemic humility. A failed prediction does not automatically identify which sentence is false.

The bad lesson would be that no claim can be tested. Scientific practice responds by designing tests that vary or stabilize different parts of the network. Independent instruments can challenge calibration assumptions. Different initial conditions can test a law across regimes. Replication teams can use the same data with alternative pipelines. Interventions can isolate causal components. A theory can therefore be tested holistically, but not indiscriminately.

Empiricist insight	What it protects	What it cannot by itself supply
Public evidence	Shared access and correction	A theory-neutral observation language
Formal clarity	Explicit inferential commitments	A choice of relevant variables
Probability	Graded belief and uncertainty	A unique prior or causal interpretation
Operational connection	Contact between concepts and procedures	A guarantee that the operation measures the intended target

Table 2.1: What the logical empiricist tradition contributes, and where it needs supplementation.

2.6 What survives

The logical empiricist tradition contributes four durable norms.

First, scientific claims should be connected to public procedures. A claim that cannot alter any observation, intervention, or inferential practice is not a scientific claim merely because it uses technical words.

Second, the relation between theory and evidence should be made explicit. Scientists should state what assumptions a result depends on and what would change if those assumptions were relaxed.

Third, uncertainty should be represented rather than hidden behind categorical language. A point estimate without an error model can be less informative than a range with clear limitations.

Fourth, philosophical arguments benefit from formal reconstruction. Logic and probability do not replace empirical inquiry, but they expose ambiguities.

ties and show exactly where an inference requires an extra premise.

What does not survive is the idea of a purely given foundation, a universal translation into observation language, or an automatic rule for induction. Objectivity will require a broader account of constraints: experimental, statistical, causal, historical, and social.

Chapter summary

Logical empiricism is best understood as a diverse attempt to connect meaning and justification to public evidence and logical structure. Its strongest achievements are the demand for inspectable procedures, explicit inferential commitments, probabilistic uncertainty, and clarity about theoretical terms. Its weaknesses arise when observation is treated as theory-neutral or when confirmation is expected to yield certainty. Holism shows that tests involve networks of assumptions, but independent instruments, interventions, and replications can still distribute and localize criticism.

Objections and discussion.

1. Does a claim become scientific when it has a possible observation, or does it also need a reliable measurement path?
2. Should a theoretical entity be accepted if it improves prediction but does not support intervention?
3. How should two investigators with different priors resolve a disagreement?

Further reading. Read Reichenbach (1938) for a mature probabilistic empiricism, Hempel (1965) for the logic of explanation, Quine (1951) for holism, and the historical overview of logical empiricism in the relevant essays collected by Carnap and his contemporaries.

Falsification, Revolutions, and Research Programmes

Chapter guide. This chapter moves from the logic of evidence to the history of scientific change. Popper's falsificationism, Kuhn's account of normal and revolutionary science, Lakatos's research programmes, and Feyerabend's methodological pluralism are treated as complementary diagnoses of different pressures within science. The central question is whether historical change destroys rational progress or reveals a more complex form of it.

The logical empiricist image of science was often read as a picture of hypotheses confronting observations one by one. Karl Popper objected that this picture misunderstands the asymmetry between confirmation and criticism. Universal statements cannot be conclusively verified by finite observations, but a universal statement can conflict with a particular observation once the auxiliary conditions are fixed. Popper therefore emphasized conjectures, deductive consequences, and attempts at refutation (Popper, 1959).

Later historians and philosophers showed that scientific practice does not fit a simple rule of immediate abandonment. Anomalies are common. Instruments fail. A theory may be central to a research programme and worth preserving while an auxiliary assumption is revised. Scientific communities work within inherited standards and problem-solutions. The important task is not to choose between logic and history but to understand how criticism operates inside historical practice.

3.1 Popper's strongest falsificationism

Popper's positive thesis is not that scientists can mechanically prove theories false. A scientific theory is exposed to risk when it rules out possible observations and when its predictions are precise enough to be vulnerable.

A theory that accommodates every possible outcome by changing its interpretation after the fact is not thereby scientific.

This is a criterion of *testability*, not a complete description of scientific careers. It explains why a bold theory can be epistemically valuable even before it is established: it tells us what the world would have to be like for the theory to fail. It also explains why ad hoc adjustments are suspicious when they merely protect a claim from evidence without generating new constraints.

Popper's account remains useful for demarcation and scientific honesty. It is not enough to collect confirming examples. A serious claim should expose itself to possible failure. But two qualifications are necessary.

First, observations do not test isolated theories. A failed prediction can reflect an instrument error, a hidden confounder, or a mistaken initial condition. Second, a theory may rationally survive an anomaly when it is part of a productive programme whose other consequences are successful. Immediate falsification is too crude as a rule of practice.

3.2 Kuhn and the structure of scientific change

Thomas Kuhn's account of scientific revolutions begins with a historical observation: much of mature science is not a continuous search for decisive tests of foundational principles. Scientists work within a paradigm, solving puzzles using shared exemplars, standards, instruments, and concepts (Kuhn, 1962). This normal science is conservative in a productive sense. It concentrates effort and permits detailed work that would be impossible if every assumption were reopened at every moment.

A paradigm can nevertheless generate persistent anomalies. When the anomalies interact with practical failures, new instruments, or a rival conceptual framework, a period of crisis may arise. A revolution changes not only an isolated theory but the questions that count as important, the objects that are salient, and the standards by which solutions are judged.

The strongest reading of Kuhn is not that scientists live in private worlds with no contact across paradigms. Kuhn explicitly retained a notion of progress in puzzle-solving and argued that later theories often extend the range and precision of problems that can be addressed. His discussion of incommensurability concerns difficulties of translation and comparison, not a claim that all theories are equally good.

Kuhn's insight is that standards are partly embodied in practice. A criterion

such as simplicity or accuracy does not decide theory choice without a background understanding of what counts as a relevant problem and a workable representation. The danger is that this insight can slide into relativism if the shared world disappears from the account. Anomalies are not merely failures to fit a community's expectations; they can be persistent effects generated by the world.

3.3 Lakatos and the time-scale of criticism

Imre Lakatos attempted to retain Popperian rationality while acknowledging the historical endurance of research traditions. A research programme has a “hard core” of central commitments, a “protective belt” of auxiliary hypotheses, and heuristic rules guiding the development of the programme (Lakatos, 1970; 1978). Scientists do not abandon a programme after every anomaly. They compare programmes over time.

A programme is progressive when modifications produce new empirical content and some of that content is corroborated. It is degenerating when changes primarily accommodate known facts without leading to novel success. The standard is comparative and historical. A theory is not condemned by a single bad result, but a programme can lose rational support if a rival repeatedly generates and confirms more demanding results.

Lakatos improves both Popper and Kuhn. He captures the importance of research continuity without treating persistence as irrational, and he turns the vague idea of paradigm success into a comparative methodological question. But the criterion is not mechanically decisive. The distinction between a novel prediction and a post hoc accommodation can depend on how a research history is reconstructed. A degenerating programme can recover after a long pause. Science also contains multiple programmes that address different targets and cannot be ranked on a single scale.

3.4 Feyerabend and the criticism of methodological monism

Paul Feyerabend argued that no fixed methodological rule captures the history of major scientific advances. The case of Galileo, for example, involved conceptual, instrumental, rhetorical, and experimental strategies that would look illegitimate if judged by a narrow rule of confirmation. Feyerabend's provocation was directed not only at method but at the political authority that can arise when one method is treated as a monopoly (Feyerabend, 1975).

The charitable interpretation is pluralist rather than nihilist. Scientific progress can require changing the standards by which a problem is approached. Alternative theories can reveal the hidden assumptions of a dominant theory. A period of methodological diversity may be epistemically useful because it prevents a community from becoming blind to its own categories.

The strongest objection to Feyerabend is that “anything goes” cannot itself be a sufficient norm. Some practices are more transparent, more reliable, and more capable of correcting error than others. A community that suppresses criticism or protects a claim from every possible test is not rescued by historical pluralism. RCR accepts methodological diversity but asks what constraints the method creates and whether those constraints are genuinely world-sensitive.

3.5 Are revolutions rational?

The phrase “scientific revolution” can suggest a discontinuity so complete that rational comparison becomes impossible. There is a real issue here: concepts change their meaning when their inferential role changes. The term “mass” in Newtonian mechanics is not simply the same term as relativistic mass. Some standards shift with the problem situation.

Yet complete discontinuity is too strong. A later theory may preserve limiting relations, numerical results, experimental techniques, or causal interventions from an earlier one. Old concepts may survive as approximations within a restricted domain. Even when vocabularies change, laboratories, bodies, detectors, and failed predictions can constrain both sides.

Historical continuity principle. Rational comparison across theories does not require a vocabulary-free standpoint. It requires enough shared constraints that rival representations can be linked to common phenomena, operations, or consequences, even when their concepts are not fully intertranslatable.

This principle also clarifies why scientific progress is not mere accumulation. A later theory can improve a representation by changing the dimensions along which the phenomenon is understood. Progress may be conceptual, predictive, causal, experimental, or organizational. It is not a single scalar quantity.

3.6 The lesson for objectivity

These traditions yield a layered picture of scientific rationality.

Tradition	Key achievement	Required correction
Popper	Risk, criticism, testability	Tests involve auxiliary assumptions and time
Kuhn	Practice, paradigms, conceptual change	Revolutions remain constrained by shared phenomena
Lakatos	Comparative research programmes	Progress is plural and historically underdetermined
Feyerabend	Methodological pluralism and criticism of monopoly	Pluralism needs standards of transparency and error correction

Table 3.1: Four accounts of scientific change and the constraints needed to make them compatible.

Scientific objectivity cannot be identified with resistance to all change. A dogmatic theory can survive anomalies. Nor can it be identified with change itself. A community can shift because of fashion, power, or institutional pressure. Objectivity concerns the quality of the constraints that govern change: whether anomalies are investigated, alternatives are developed, evidence is independently checked, and revisions increase the range and reliability of world-directed inference.

Chapter summary

Popper shows why risky claims and attempts at refutation matter. Kuhn shows that scientists work within paradigms and that revolutions change problems and standards. Lakatos explains why rational criticism often operates over the history of research programmes rather than through instant rejection. Feyerabend warns against methodological monopoly and reminds us that new methods can be necessary for new problems. Together they support a fallibilist, historical, and plural account of rationality, provided that theory change remains constrained by common phenomena, interventions, and error correction.

Objections and discussion.

1. When should an anomaly be treated as evidence against a theory rather than as a reason to revise an auxiliary assumption?
2. Can a research programme be progressive in one domain and degenerating in another?
3. Does conceptual incommensurability prevent rational comparison, or only make comparison more demanding?

Further reading. Compare Popper (1959), Kuhn (1962), Lakatos (1970), and Feyerabend (1975) directly. For critical discussion of historicist rationality, consult the relevant scholarship on scientific change and the essays collected in Lakatos and Musgrave (1970).

Realism, Anti-Realism, and the Problem of Success

Chapter guide. This chapter asks what the success of science licenses us to believe. It examines constructive empiricism, scientific realism, structural realism, entity realism, pragmatism, inference to the best explanation, and the pessimistic meta-induction. The chapter's conclusion is a layered conception of truth and a graded account of ontological commitment.

The central realist argument is simple. Mature scientific theories often predict novel phenomena, integrate diverse observations, guide new instruments, and support interventions. It would be a remarkable coincidence if the theoretical entities and structures responsible for this success were entirely fictional. Scientific realism therefore recommends belief in at least some unobservable aspects of the world described by our best theories (Chakravartty, 2007).

The anti-realist replies that success need not reveal truth. A theory can be empirically adequate: it can correctly describe what is observable while remaining agnostic about unobservable entities. Bas van Fraassen's constructive empiricism is not a claim that electrons do not exist. It is a claim about the epistemic attitude required by science: acceptance may involve belief that a theory is empirically adequate without belief that its unobservable posits are true (van Fraassen, 1980).

Both positions contain an important distinction. The question is not whether a scientist may use a model containing unobservables. Scientists plainly do. The question is what attitude the use and success of the model rationally warrant.

4.1 Constructive empiricism

Constructive empiricism avoids a crude instrumentalism. It can accept that scientific theories are true-or-false claims, that they explain phenomena, and that scientists can rationally adopt them for practical purposes. Its distinctive requirement is epistemic modesty: belief need extend only to observable adequacy.

This view has several strengths. It respects the historical fact that successful theories are often replaced. It avoids treating every useful mathematical posit as an entity. It separates the aim of science from the psychology of individual scientists. And it recognizes that seeing through a microscope or detector involves a complex notion of observability rather than an absolute line between the visible and invisible.

Its difficulty is that the observable/unobservable distinction can be unstable and method-dependent. A virus, a distant galaxy, and a magnetic field do not form a natural epistemic class. More importantly, empirical adequacy can be achieved by models whose internal structure supports counterfactual and interventional reasoning. If a theory's unobservable components are repeatedly used to make new kinds of contact with the world, permanent agnosticism may be too restrictive.

4.2 Scientific realism and the no-miracles intuition

Realists often appeal to the no-miracles argument: the ability of science to predict and intervene would be extraordinary if theories were not at least approximately true. The argument is not a deductive proof. It is an inference to the best explanation (IBE). We compare explanations of scientific success and regard truth-tracking as one candidate.

Gilbert Harman used the phrase “inference to the best explanation” for the pattern in which a hypothesis is selected because it would, if true, explain the evidence better than alternatives ([Harman, 1965](#)). IBE is not a license to accept whatever story feels satisfying. A good explanatory hypothesis should have independently supported consequences, expose itself to failure, and outperform rivals in a specified domain.

The no-miracles argument is strongest when success is robust across independent routes: a theory predicts an effect before it is observed, survives instrumental changes, connects different phenomena, and supports interventions that would not be possible if the relevant structure were absent. It is weaker when the success is merely a fit to existing data, depends on

flexible parameters, or is achieved by many unrelated models.

4.3 The pessimistic meta-induction

The historical record supplies an obvious challenge. Many past theories were successful and yet contained entities or principles that later scientists rejected. The luminiferous ether supported a useful wave theory of light. Phlogiston organized chemical phenomena before oxygen chemistry. Classical mechanics remains reliable in many domains although its ontology is not the final physical description.

The pessimistic meta-induction argues that our current theories may be like those past theories: successful for a time but false in their unobservable commitments. This is not a proof of anti-realism. It is a warning against treating present success as a guarantee of total truth.

The realist response often distinguishes between abandoned structures and retained structures. A later theory may reject an entity while preserving equations, symmetry relations, limiting cases, or causal dependencies. John Worrall's structural realism emphasizes precisely this possibility (Worrall, 1989). The theory changes, but some relations survive.

The response is plausible but incomplete. Mathematical structure is not a single object. One must specify which structure is retained, at what level, and under what transformations. A merely formal similarity can survive while the physical interpretation changes. RCR therefore treats structural continuity as evidence, not as an automatic guarantee of truth.

4.4 Structural realism

Structural realism comes in epistemic and ontic forms. Epistemic structural realism says that science gives us knowledge of relations or structure even if knowledge of the intrinsic nature of entities remains limited. Ontic structural realism makes the stronger metaphysical claim that structure is more fundamental than objects.

The epistemic version fits an important pattern in theory change. A relation such as a mathematical dependency may survive when the entities used to interpret it change. The relational pattern can therefore be a better target of realist commitment than the full ontology of a theory.

The ontic version is more ambitious but faces questions about what a structure is without relata, how structures are individuated, and how causal powers fit into a purely relational picture. RCR remains neutral between object-first and structure-first metaphysics at the ultimate level.

It claims only that scientific evidence often supports structured constraints more securely than an exhaustive inventory of intrinsic natures.

4.5 Entity realism and intervention

Ian Hacking shifted attention from representation to intervention. When scientists can use an entity to create stable effects in other systems, there is a strong practical reason to believe that the entity is not merely a calculational fiction (Hacking, 1983). This is especially persuasive in experimental settings where the entity is not just inferred from the same data that the theory was designed to fit.

Intervention is not infallible. A successful manipulation can be described by different ontologies. An instrument may be responding to a field, a mechanism, or a composite process. Experimental control can also be achieved through a model whose components are only approximately isolated.

The lesson is therefore conditional. Intervention raises the ontological status of a posit when the effect is robust under changes in apparatus, calibration, and background assumptions, and when the posit makes a difference to outcomes that were not used to define it. This will become one of the criteria in RCR's commitment ladder.

4.6 Pragmatism without relativism

Pragmatist traditions, especially Peirce's account of inquiry, emphasize that belief is a habit of action and that the meaning of an idea lies partly in its conceivable practical consequences (Peirce, 1877–1878). This is not the claim that truth means whatever works for a person. Peirce's ideal of inquiry is public, fallible, and oriented toward a community that could continue testing the claim.

Pragmatism corrects two realist errors. First, it reminds us that theories are tools for asking questions and changing situations, not museum labels attached to a finished world. Second, it prevents the search for truth from becoming detached from the practices that give concepts their use.

It also needs correction. Immediate usefulness can reward a false model, especially when a short-term prediction is obtained through a fragile correlation. RCR distinguishes pragmatic adequacy from correspondence and causal adequacy. A model can be useful while its ontological interpretation remains uncertain.

4.7 A layered account of truth

The debate becomes less confusing when different kinds of scientific success are separated.

Layer	Question	Typical evidence
Correspondence	Does the claim identify what exists or occurs?	Independent detection, historical trace, convergent measurement
Structural accuracy	Does the representation preserve relevant relations?	Invariance, limiting cases, cross-domain integration
Interventional adequacy	Does changing the represented variable change outcomes as expected?	Experiment, manipulation, mechanism, counterfactual test
Pragmatic adequacy	Does the representation support reliable inquiry or action?	Forecasting, control, diagnosis, design

Table 4.1: Four layers of scientific success. They support one another but do not collapse into one another.

A model may be structurally accurate in a restricted regime while being literally false outside it. A variable may be useful for intervention without being a fundamental entity. A theory may correspond to a real mechanism while its mathematical idealization misstates the mechanism's fine detail. The proper question is not whether the whole theory is simply true, but which claims, structures, and interventions have earned which kinds of commitment.

Layered truth principle. Scientific truth is not one undifferentiated property. A claim may be correct about a phenomenon, structurally correct about a relation, causally adequate for an intervention, and pragmatically useful to different degrees. Objectivity improves when these layers are reported separately rather than compressed into one label.

4.8 What realism should commit us to

RCR rejects both the demand for literal truth of every model component and the refusal to believe in any unobservable. It recommends graded

commitment.

1. Believe an observable event when its measurement path is reliable and alternative explanations have been tested.
2. Treat a structure as objectively supported when it remains stable across independent representations and transformations.
3. Treat an entity or mechanism as strongly supported when intervention, convergence, and counterfactual robustness all point to it.
4. Treat a theoretical ontology as provisional when its success is confined to prediction or depends on highly flexible idealization.

This is not an attempt to turn realism into a numerical score. It is a discipline of separating what the evidence supports from what the imagination adds. The final commitment should be proportional to the strongest constraint that the inquiry has actually established.

Chapter summary

Constructive empiricism protects epistemic modesty; scientific realism explains why robust success can support belief in unobservables; structural realism highlights continuity of relations; entity realism emphasizes intervention; and pragmatism links meaning to public inquiry and action. The pessimistic meta-induction warns that successful theories can contain false commitments. The proposed response is layered truth and graded commitment: correspondence, structural, interventional, and pragmatic adequacy are distinct dimensions of scientific success.

Objections and discussion.

1. Is an entity more credible when it can be manipulated, or can intervention itself be theory-dependent?
2. Which parts of a modern scientific theory seem most secure: entities, equations, structures, mechanisms, or measurement procedures?
3. Can empirical adequacy be enough for practical science even when ontology is unsettled?

Further reading. Read van Fraassen (1980), Hacking (1983), Worrall (1989), Fine (1984), Chakravartty (2007), and Stanford (2006)

as contrasting positions in the realism debate.

Observation, Measurement, and the Making of Data

Chapter guide. This chapter examines how observations become data. It argues that theory ladenness does not make observation arbitrary. Instead, objectivity depends on a traceable measurement path, calibrated instruments, controlled sources of error, and the convergence of distinct ways of making the target answer.

A datum is not a small piece of reality waiting to be collected. It is the result of an interaction between a target, an instrument, a procedure, a record, and an interpretive model. The fact that data are made does not mean that they are made up. A footprint is produced by a foot and a surface, but the scientific significance of the footprint depends on a method of reading it.

This point matters because scientific disputes often focus on whether an observation is theory-laden. The correct answer is yes: observations rely on concepts and instruments. But the conclusion that observations are therefore subjective does not follow. A map is concept-laden and can still be constrained by terrain. The relevant question is whether the concepts and instruments are arranged so that the result can resist convenient reinterpretation.

5.1 Theory-ladenness without observational chaos

Norwood Hanson and later philosophers emphasized that what a scientist sees depends on training and expectation. A radiologist and a novice may look at the same image and notice different patterns. A physicist sees

a detector event as evidence of a particle only within a theoretical and instrumental context. Wilfrid Sellars argued that the “given” is not a self-interpreting foundation of knowledge (Sellars, 1963).

The strongest lesson is semantic and inferential. There is no observation sentence whose meaning is entirely independent of all concepts. The novice and the expert do not occupy identical epistemic positions. But the existence of different interpretations also creates a testable difference. An expert’s interpretation can be assessed by calibration, successful prediction, independent readings, and the ability to distinguish real signals from artifacts.

Theory-ladenness has two directions. Theory guides attention and helps design instruments. Evidence also disciplines theory by exposing predictions that fail, measurements that do not replicate, and interventions that produce unexpected effects. The relation is therefore reciprocal rather than one-way.

5.2 The measurement path

To evaluate a measurement, RCR asks for a path from target to claim:

1. **Target:** What property, event, or process is being measured?
2. **Interaction:** How does the target affect the instrument or observer?
3. **Calibration:** How is the instrument’s response related to a standard or independently measured quantity?
4. **Transformation:** What mathematical or computational operations turn the raw signal into a reported value?
5. **Error model:** Which systematic and random errors could distort the result?
6. **Interpretation:** What claim is the result taken to support, and what background assumptions connect it to the target?

This path is not a demand that every measurement be philosophically reconstructed from scratch. It is a practical test of where objectivity could fail. A value can be precise but invalid if it measures the wrong construct. A survey can be representative but biased by nonresponse. A detector can be sensitive but unable to distinguish signal from background. A psychological scale can be reliable across repeated administrations while failing to measure the concept its label suggests.




Figure 5.1: A measurement path. Every arrow is a potential site of error and also a site where independent validation can be added.

5.3 Operationalism and construct validity

Operationalism insists that a concept should be defined by an operation. The view has a useful corrective: a definition that cannot guide any procedure is not yet an empirical concept. But one concept may have several legitimate operations. Temperature can be measured through expansion, resistance, radiation, or thermodynamic relations. Intelligence, poverty, and social trust also admit multiple operationalizations.

An operation is therefore not identical to a definition. It is an instrument for producing evidence about a construct under conditions. Convergent measurement asks whether different operations agree when they should. Discriminant measurement asks whether they remain distinct from neighboring constructs. Known-group tests ask whether a measure behaves differently in groups that the theory says should differ. These tests do not eliminate interpretation, but they distribute it across independent routes.

5.4 From data to phenomena

James Bogen and James Woodward distinguished data from phenomena (Bogen and Woodward, 1988). Data are local, often noisy records: a set of instrument readings, images, or observations. Phenomena are stable patterns that data reveal through analysis, replication, and background knowledge. The distinction explains why a single observation can be weak evidence while a robust phenomenon can support a theory.

For example, a particular climate station reading may be affected by instrument placement, urbanization, or missing values. A long-run warming pattern is not simply one more reading; it is a phenomenon supported by many records, corrections, physical mechanisms, and independent indicators. A single detector event may be ambiguous, while a distribution of events with a predicted energy relation is a phenomenon.

The distinction also protects against two errors. Data are not automati-

cally theory-free facts, but phenomena are not inventions of a theory. A phenomenon is a stabilized pattern that remains identifiable when the routes to it are varied and criticism is allowed to operate.

5.5 Error, uncertainty, and precision

Precision concerns the spread of repeated measurements. Accuracy concerns closeness to the target value. A measurement can be precise but systematically wrong. A model can estimate a tiny standard error because it ignores a major source of uncertainty. A survey can report many decimal places while its conceptual validity is poor.

RCR distinguishes three forms of uncertainty.

- **Aleatory uncertainty** concerns variation in a process or outcome that is treated as stochastic for the purpose of analysis.
- **Epistemic uncertainty** concerns limited knowledge of parameters, models, initial conditions, or measurement error.
- **Scope uncertainty** concerns whether a result transfers to a different population, scale, time, or institutional setting.

The distinction is not merely terminological. More observations may reduce sampling uncertainty but do little to correct a biased measurement. A larger sample may make a misleading operational definition more precise. A model ensemble may explore parameter uncertainty while leaving structural uncertainty untouched.

5.6 Instrumental realism

Ian Hacking's emphasis on experimentation suggests a useful form of realism: instruments can create a stable relation between a target and an effect. But instrumental success must be unpacked. A detector's output can be used to control a beam, diagnose a condition, or produce a new phenomenon. The stronger the independence between the instrument's construction and the claim being tested, the stronger the resulting support.

One should ask:

1. Can the result be reproduced with a different instrument principle?
2. Can the calibration be checked against an independent standard?
3. Does the target support interventions that were not used to define it?
4. Do failed attempts reveal a boundary rather than disappear through post hoc interpretation?

This standard is demanding but not impossible. It is common in metrology, particle physics, medicine, and engineering. In social research, the analogous tests may involve alternative measures, preregistered coding, blinded assessment, administrative records, or natural experiments.

5.7 Observation in the social sciences

Social categories present a special issue because measurement can change the thing measured. A census category can become an administrative identity. A test score can alter self-perception. A policy indicator can become a target that organizations optimize. This does not make the category unreal, but it means that the measurement relation is reflexive.

Reflexive measurement principle. When the act of classification can change behaviour, the measurement path must include feedback from the classification to the target. Objectivity requires reporting that feedback, not pretending that the target is static.

Reflexive measurement is one reason why social science needs both causal models and institutional analysis. A variable can be a real social fact because people and organizations act on it, while its meaning and effects depend on rules that can change.

Chapter summary

Observation is conceptually mediated but not arbitrary. A measurement becomes objective to the extent that its path from target to reported result is traceable, calibrated, error-aware, and open to independent checks. Data are local records; phenomena are stabilized patterns that survive analysis and replication. Precision does not guarantee validity, and social measurements require attention to feedback. Theory and evidence constrain one another through carefully designed measurement paths.

Objections and discussion.

1. Which is more dangerous in a measurement: random noise, systematic bias, or a construct that is poorly defined?
2. Can two measurements of the same construct legitimately disagree?
3. How should a researcher report a result when the target changes in response to being measured?

Further reading. See Hacking (1983) on experiment and intervention, Bogen and Woodward (1988) on data and phenomena, Sellars (1963) on the critique of the given, and Frigg and Hartmann (2006) for models and measurement in scientific practice.

Models, Explanation, Prediction, and Laws

Chapter guide. Models are treated as selective mediators between theory and world. The chapter explains how idealized or literally false models can be scientifically useful, compares major accounts of explanation, and asks what laws of nature are doing in actual scientific practice. It then connects reduction and emergence to the idea of scale-specific constraints.

Scientific theories are rarely used as complete descriptions. Scientists build models: simplified systems, equations, diagrams, simulations, analogies, mechanistic sketches, statistical representations, and idealized populations. A model may contain frictionless planes, perfectly rational agents, isolated genes, representative households, point particles, or infinite populations. None of these may exist exactly. Yet the models can reveal dependencies, generate predictions, and guide interventions.

The philosophical mistake is to treat usefulness and literal truth as competitors. The right question is what a model represents, under which transformations, and for which purpose.

6.1 Models as mediators

Mary Morgan and Margaret Morrison describe models as mediators that are neither simply derived from theory nor copied from the world ([Morgan and Morrison, 1999](#)). A model can borrow resources from several theories, data sources, and experimental practices. It provides a manipulable object for reasoning.

Ronald Giere's account is similarly selective. A model represents a real system to a degree and for a purpose when the relevant similarities are identified ([Giere, 1988; 2004](#)). A model of a population may represent its age distribution but not its cultural history. A model of a pendulum may

represent period under small-angle conditions while ignoring air resistance. The model's success is conditional on what is preserved.

A representation relation is therefore not binary. We can write informally:

$$M \xrightarrow[\text{purpose, scale, conditions}]{\text{represents}} T,$$

where M is a model and T is a target system. The arrow is labelled because the same model can represent different features of a target under different purposes. The representation is strong when relevant relations are preserved and weak when the model changes the target's behaviour in the dimension that matters.

6.2 Idealization and abstraction

An idealization deliberately removes factors. Abstraction selects a feature while leaving other features unspecified. These are not the same. A frictionless plane removes friction by setting a parameter to zero. A map abstracts from buildings when it preserves roads and elevation. An agent-based model may abstract from individual biographies while idealizing agents as following a simple rule.

A useful idealization should satisfy three conditions.

1. **Relevance:** the removed feature should not be central to the question under investigation.
2. **Stability:** the result should remain approximately valid when the removed feature is reintroduced within the intended range.
3. **Boundary:** the model should state where the idealization breaks down.

A model can be literally false and structurally accurate. For a simple example, consider a population model with exponential growth:

$$N(t) = N_0 e^{rt}.$$

Here $N(t)$ is the population at time t , N_0 is the initial population, and r is a constant net growth rate. No real population grows with an exactly constant r forever. The model can still represent a local growth regime, estimate a rate, or provide a baseline against which more complex models are compared.

The danger is not idealization itself. The danger is forgetting that a result about the model is being treated as a result about the target without a

transport argument.

6.3 Prediction and explanation

Prediction and explanation overlap but are not identical. A model can predict an outcome using a reliable correlation without revealing why the outcome occurs. A mechanism can explain a process without providing a precise long-run forecast when initial conditions are unknown or the system is chaotic.

The deductive-nomological model associated with Hempel and Oppenheim treats an explanation as a deduction of the explanandum from laws and initial conditions (Hempel and Oppenheim, 1948; Hempel, 1965). Its great virtue is clarity: it asks which premises entail the result. Its limitation is that logical derivation alone does not guarantee explanatory relevance. An irrelevant law can be included in a valid deduction, and probabilistic explanations do not entail their outcomes.

Unificationist approaches evaluate explanations by how many phenomena they derive from a small set of argument patterns. This captures a real value: compressing unrelated facts into one framework can reveal a shared structure. But simplicity can mislead when the world is heterogeneous.

Causal and mechanistic approaches focus on how a phenomenon is produced. Salmon emphasized causal structure and statistical relevance (Salmon, 1984). Mechanistic philosophers analyze organized parts and activities whose interaction produces a behavior (Bechtel and Abrahamsen, 2005; Craver, 2007). These accounts explain why an intervention or process makes a difference, not only why an equation fits.

Explanatory adequacy principle. A scientific explanation is adequate only relative to a question. For a how-question, a mechanism may be required. For a prediction question, a calibrated statistical model may be enough. For a why-question about a general pattern, unification or a law-like relation may matter. No single form of explanation is the universal measure of science.

6.4 Laws of nature and local regularities

A law is often imagined as a universal sentence that describes what nature does in every circumstance. Actual science is more complicated. Laws may be exact ideal relations, approximate regularities, symmetry principles, constraints on possibilities, or statements that hold only under a *ceteris paribus* clause.

Nancy Cartwright argues that fundamental laws can be highly idealized and that scientific explanation often proceeds through models that construct a “simulacrum” of a phenomenon rather than by direct deduction from universal laws (Cartwright, 1983; 1999). This is not a denial of reality. It is a warning that a law’s role in a model may differ from a law’s metaphysical status.

D. M. Armstrong’s account treats laws as relations among universals (Armstrong, 1983). Other accounts treat them as regularities, dispositions, or modal constraints. RCR does not settle the ultimate metaphysics by definition. It says that a law-like claim earns scientific authority when it captures a stable constraint across a specified range of interventions and conditions, and when deviations can be explained by additional factors rather than ignored.

6.5 Reduction and emergence

Reduction has at least three meanings:

1. **Ontological reduction:** higher-level entities are nothing over and above lower-level entities.
2. **Explanatory reduction:** higher-level laws or explanations can be derived from lower-level laws.
3. **Methodological reduction:** a problem is best investigated by isolating lower-level components.

These meanings can come apart. A social institution may be constituted by people and rules while still requiring concepts that are not useful in particle physics. A biological pathway may be physically realized while explanations at the molecular, cellular, and organismic levels remain distinct. Nagel’s classical model of reduction emphasizes derivation and bridge principles (Nagel, 1961); later work has shown that reduction in practice can be local, approximate, and model-dependent.

Emergence is not a magic word for ignorance. A property is emergent in a scientifically meaningful sense when it depends on lower-level organization but has stable patterns, explanatory roles, or causal powers at a higher level. Phase transitions provide cases where limiting behavior and scale can make higher-level regularities more informative than a microscopic description alone (Batterman, 2002).

RCR treats levels as nested constraint regimes. A higher-level model is not automatically anti-realist because it cannot be replaced by a lower-level description in practice. Nor is every higher-level regularity fundamental.

The question is whether the higher-level constraint is stable, manipulable, and explanatory within its domain.

6.6 Simulation and computational models

Computer simulations are models executed as algorithms. They can explore systems for which analytic solutions are unavailable, but their results are not automatically empirical evidence. A simulation inherits assumptions from its equations, discretization, parameter values, boundary conditions, and code. It gains epistemic strength when these components are validated against analytical results, historical data, experiments, or independent simulations.

A simulation can also discover emergent behavior. The discovery is not that the code produced a surprise in a metaphysical vacuum. It is that a specified set of rules generates a pattern that was not obvious from inspection and that can be compared with a target system. The comparison requires a bridge between the computational state and the phenomenon.

6.7 The model's honesty condition

Every scientific model should make four things visible:

1. what it treats as the target;
2. what it idealizes or leaves out;
3. which output is a prediction, explanation, or exploratory possibility;
4. how the output was validated and where it should not be transported.

This is the *model's honesty condition*. A model need not contain every detail to be honest. It needs to distinguish internal consequences from claims about the world.

Chapter summary

Models mediate between theory and world. They are selective, purpose-relative, and often idealized, but that does not make them arbitrary. A model is scientifically valuable when it preserves relevant structure, remains stable under appropriate changes, and states its boundary conditions. Prediction, explanation, and intervention can come apart. Laws operate through idealized models and local regimes. Reduction and emergence are best treated as relations among levels of constraint rather than as all-or-nothing doctrines.

Objections and discussion.

1. When does an idealization become a distortion rather than a useful abstraction?
2. Can a model explain a phenomenon if it cannot predict it?
3. Does a higher-level explanation remain objective when a lower-level description is physically complete?

Further reading. See Morgan and Morrison (1999), Giere (1988, 2004), Cartwright (1983, 1999), Frigg and Hartmann (2006), Salmon (1984), and Craver (2007).

Causality, Mechanisms, and Complexity

Chapter guide. This chapter distinguishes association, intervention, and explanation. It introduces potential outcomes, directed acyclic graphs, structural causal models, interventionist explanation, mechanisms, and transportability. The central claim is that causal knowledge is possible when the inquiry creates or locates stable constraints on change, not merely patterns of co-occurrence.

Causal claims answer a practical question: what would happen if something were changed? A correlation describes how variables vary together. A causal claim says that a difference in one variable would produce a difference in another under specified conditions. The distinction is not metaphysical decoration. Medical treatment, economic policy, education, and climate intervention all depend on it.

Causality is difficult because the same association can arise from direct causation, common causes, selection, measurement error, feedback, or chance. No statistical method can identify a causal interpretation without assumptions about design, time, variables, and the process generating the data. But the presence of assumptions does not make causal knowledge impossible. It means that causal claims should expose the assumptions that make them testable.

7.1 Potential outcomes and the missing counterfactual

The potential-outcomes framework imagines that each unit has an outcome under each possible treatment. Let $Y_i(1)$ be the outcome for unit i if it receives treatment and $Y_i(0)$ the outcome if it does not. The individual causal effect is

$$\tau_i = Y_i(1) - Y_i(0).$$

For a given unit, we observe only one of these two potential outcomes. This is the fundamental problem of causal inference (Rubin, 1974; Holland, 1986). A randomized experiment solves part of the problem by making treatment assignment independent of potential outcomes in expectation. The average difference between treatment and control groups can then estimate an average causal effect, provided the treatment and outcome are defined and measured properly.

Randomization is not a universal substitute for thought. It may have limited external validity, noncompliance, attrition, spillovers, or ethical constraints. A trial can estimate what happened under its treatment protocol without establishing what would happen under a different dose, population, or institution.

7.2 Graphs and confounding

A directed acyclic graph (DAG) represents assumed causal direction. In a simple confounding structure, C affects both treatment T and outcome Y :




Figure 7.1: A confounding graph: C is a common cause of treatment T and outcome Y . The association between T and Y can mix the treatment effect with the path through C .

If C is measured and the graph is credible, conditioning on C may block the back-door path from T to Y . If C is unmeasured, researchers may need randomization, an instrumental variable, a natural experiment, negative controls, sensitivity analysis, or a different design. The graph does not prove the causal story. It makes the story visible enough to argue about.

Judea Pearl’s structural causal models distinguish observation from intervention. Informally, $\mathbb{P}(Y \mid T = t)$ describes the outcome among cases observed with $T = t$, whereas $\mathbb{P}(Y \mid \text{do}(T = t))$ describes what would happen if T were set to t by an intervention that removes the usual causes

of treatment (Pearl, 2009). The distinction explains why an association can fail to answer a policy question.

7.3 Bayesian reasoning about causal claims

Bayesian reasoning can represent uncertainty about causal hypotheses. Suppose H_1 says a treatment has a positive effect and H_0 says it has no effect. The evidence E changes their odds according to

$$\frac{\mathbb{P}(H_1 | E)}{\mathbb{P}(H_0 | E)} = \frac{\mathbb{P}(E | H_1)}{\mathbb{P}(E | H_0)} \times \frac{\mathbb{P}(H_1)}{\mathbb{P}(H_0)}.$$

The first ratio is the likelihood ratio: how much more expected the evidence is under H_1 than H_0 . The second ratio expresses prior odds. A strong randomized design can make the likelihood ratio informative by reducing confounding. A weak observational design may produce a large apparent effect while leaving several causal models compatible with the data.

The equation does not remove causal assumptions. It shows where they enter. A prior may encode background knowledge about mechanisms. The likelihood depends on the measurement model, treatment assignment, missingness, and outcome definition. Good causal inference therefore combines probability with design and mechanism.

7.4 Interventionist explanation

James Woodward treats causal explanation in terms of stable relations under interventions (Woodward, 2003). A variable X explains a change in Y when manipulating X , while holding relevant background conditions fixed, changes Y according to a relationship that is stable enough to support counterfactual reasoning.

This account improves on correlation, but the phrase “holding fixed” hides complexity. In a biological system, manipulating one gene may alter many pathways. In an economy, changing interest rates can trigger expectations and institutional responses. In a social setting, an intervention can change the meaning of the variable itself.

RCR therefore adds an *intervention boundary*: a causal claim must state which variables are changed, which processes are allowed to respond, and which background conditions are treated as fixed. A causal law is not a universal promise that every intervention works identically. It is a scope-bound constraint on a family of interventions.

7.5 Mechanisms and causal structure

Mechanistic explanations describe parts, activities, organization, and the way their interactions produce a phenomenon. A mechanism is not merely a list of correlations. It explains how a change propagates through a system.

Mechanistic evidence is especially valuable when it connects different scales. A drug may bind to a receptor, alter a signaling pathway, change cell behavior, and improve a clinical outcome. Each link can fail independently. The mechanistic story is stronger when it is supported by intervention and measurement at several links rather than inferred from the final outcome alone.

Mechanisms are not always microscopic. A market institution, a school organization, or a legal procedure can be a mechanism when its parts and rules produce a stable social outcome. The relevant parts may be persons, norms, resources, and sanctions rather than molecules.

7.6 Complexity, feedback, and nonlinearity

Complex systems create three challenges for causal inference.

First, effects can be nonlinear. A small change may have little effect below a threshold and a large effect above it. A linear average can conceal this structure.

Second, systems can contain feedback. Treatment changes behaviour; behaviour changes treatment exposure; the outcome changes future measurement. A static DAG may need to be expanded into a time-indexed graph.

Third, many causes can interact. The effect of an intervention may depend on context, history, and network position. The average treatment effect can be scientifically useful while failing to describe any particular subgroup.

These challenges do not defeat objectivity. They make *transportability* and *heterogeneity* central. A result should be carried to a new setting only after checking which mechanisms, distributions, and institutional conditions are shared.

Causal constraint principle. A causal claim is objective to the degree that it identifies a stable change-relation supported by design, measurement, mechanism, or convergent intervention. The claim must state its intervention boundary and the conditions under which the relation is expected to travel.

7.7 What causal knowledge does not require

Causal knowledge does not require deterministic laws. An intervention can change a probability distribution without fixing an individual outcome. Nor does causal knowledge require a single metaphysical definition of cause. Different sciences may use manipulation, process transmission, counterfactual dependence, or mechanism construction as complementary modes.

Causal knowledge also does not require the elimination of all uncertainty. It requires that the uncertainty be attached to a specified causal question. A claim such as “this policy works” is too vague. A claim such as “for schools with these resources, under this implementation, assignment to this programme increases average attendance over one year by an estimated range” is criticizable and therefore epistemically useful.

Chapter summary

Correlation records association; causal inference concerns changes under intervention. Potential outcomes clarify the missing counterfactual, DAGs make confounding assumptions visible, and structural causal models distinguish observation from intervention. Mechanisms explain how effects are produced. Complexity requires attention to nonlinearity, feedback, heterogeneity, and transportability. RCR treats causal objectivity as the stabilization of change-relations under explicit intervention boundaries.

Objections and discussion.

1. Can a randomized trial establish a causal effect if treatment changes the outcome measure?
2. When is a mechanism more informative than a highly predictive model?
3. How should a causal claim be reported when its average effect hides opposite effects in subgroups?

Further reading. See Rubin (1974), Holland (1986), Pearl (2009), Woodward (2003), Hernan and Robins (2020), Salmon (1984), and Craver (2007).

Social Epistemology, Standpoints, and Values

Chapter guide. This chapter asks whether social location and institutional organization undermine or improve objectivity. It treats science as a collective achievement, examines feminist standpoint theory and situated knowledge, and distinguishes values that select questions from values that distort evidence. The proposed ideal is public, structured criticism rather than a view from nowhere.

No individual scientist can reproduce all the measurements, derivations, computer codes, field observations, and historical records on which a mature science depends. Scientific knowledge is therefore social in an unavoidable sense. The question is whether the social character is a source of objectivity, a source of bias, or both.

The answer is both. Division of labour allows specialization and independent checking. It also creates authority gradients and dependencies that can hide errors. Peer review can expose mistakes. It can also enforce conformity. Funding can support long-term inquiry. It can also redirect attention toward profitable or politically convenient questions. Social structure is not an external background to epistemology. It is part of the machinery by which evidence is produced and evaluated.

8.1 The social organization of knowledge

Social epistemology studies how people pursue truth with and through others. Alvin Goldman emphasizes testimony, expertise, and the conditions under which social processes are reliable (Goldman, 1999). A person often knows something because another person measured it, and the trustworthiness of that testimony depends on competence, incentives, access to criticism, and a record of past performance.

The division of cognitive labour can improve objectivity in at least four

ways:

1. specialists develop techniques too complex for one investigator;
2. independent groups can detect local mistakes;
3. different backgrounds make different anomalies salient;
4. public records allow later investigators to audit procedures.

But social dependence also creates vulnerabilities. A community can coordinate around a false assumption. A prestigious laboratory can make contrary evidence hard to publish. Researchers may replicate a result only when it is fashionable. A field can measure what is easy rather than what is important.

RCR therefore evaluates institutions by their *correction capacity*: the ability to expose a claim to competent, motivated, and sufficiently independent criticism. Consensus is evidence only when the process generating the consensus has a reasonable relation to truth.

8.2 Merton's norms and their limits

Robert Merton described norms associated with the scientific ethos: communal sharing of knowledge, universalistic assessment, disinterestedness, and organized skepticism (Merton, 1942). The norms are ideals, not descriptions of every laboratory. They express a valuable thought: scientific claims should be judged by reasons that are not supposed to depend on the speaker's identity or patronage.

The norms require supplementation. Universalistic assessment can become blind to how research questions were selected. Communal sharing can conflict with privacy or intellectual property. Disinterestedness does not mean that scientists have no values; it means that personal or institutional interests should not decide the evidential conclusion. Organized skepticism can be distributed unevenly, with critics demanding impossible standards from outsiders and weak standards from insiders.

8.3 Longino and transformative criticism

Helen Longino argues that objectivity is not achieved by eliminating social values from inquiry. It is strengthened by critical interaction under appropriate social conditions (Longino, 1990). The relevant conditions include recognized venues for criticism, uptake of criticism, public standards, and a degree of equality in intellectual authority.

This account turns disagreement into an epistemic resource. A criticism

is not valuable merely because it is dissenting. It must identify a reason to doubt, reveal a hidden assumption, offer an alternative, or show that a result does not travel. The community must also be able to respond rather than ignore the criticism.

Critical interaction principle. A scientific community is more objective when relevant criticism can be expressed, understood, answered, and incorporated without the critic having to possess the dominant group's institutional status.

The principle is compatible with expertise. Equality of intellectual authority does not mean that every claim has equal evidential weight. It means that access to criticism should not be determined solely by status, and that expertise itself should be accountable to public reasons.

8.4 Standpoint theory and situated knowledge

Feminist standpoint theory argues that social positions can shape what is visible and can sometimes provide epistemic advantages concerning relations of power. Sandra Harding connects standpoint theory to the idea that inquiry should begin from lives and activities that dominant institutions have treated as peripheral (Harding, 1991). Donna Haraway emphasizes situated knowledges: every perspective is partial, and objectivity requires recognizing the location and limits of the knower (Haraway, 1988).

The strongest version is not that every member of an oppressed group automatically knows the truth. Social position can generate insight, but it can also generate partiality, pressure, or exclusion. Epistemic advantage is an achievement of critical reflection and collective inquiry, not a biological property.

Standpoint insights matter because a theory can fit the experience of a dominant group while missing the mechanisms affecting others. Research on workplace safety, medical treatment, domestic labour, disability, and discrimination can improve when people who bear the consequences participate in question formation and interpretation. This is not a concession that evidence is subjective. It is an argument that the range of relevant evidence is socially organized.

8.5 Values in science

The value-free ideal is ambiguous. There is a legitimate distinction between values that determine what a result is and values that guide which questions are worth asking. Scientific work requires choices about significance, acceptable risk, measurement cost, privacy, classification, and the balance between false positives and false negatives.

Heather Douglas argues that values can have a legitimate role in decisions about acceptable error, especially where evidence informs policy (Douglas, 2009). Elizabeth Anderson shows how value judgments can guide concept formation and interpretation without making the evidence itself a matter of preference (Anderson, 2004). Max Weber distinguished empirical analysis from value relevance: values can select a problem without deciding the factual answer (Weber, 1949).

Role of value	Legitimate function	Epistemic risk
Question selection	Decide what matters to investigate	Neglect of less visible populations
Concept formation	Choose distinctions relevant to a problem	Smuggling a conclusion into a definition
Loss function	Weigh false positives and false negatives	Treating a policy preference as a fact
Evidence interpretation	Decide when uncertainty warrants action	Confirmation bias and motivated reasoning
Institutional design	Set rules for access and correction	Exclusion, capture, and authority abuse

Table 8.1: Values can enter different stages of inquiry. Their presence does not make all stages subjective, but it makes transparency necessary.

The key distinction is between *constitutive* and *evidential* values. Constitutive values help determine what to study, how to measure it, and what risks to accept. Evidential values would make a claim count as true because it serves an interest, regardless of the evidence. The first are unavoidable and can be legitimate. The second corrupt objectivity.

8.6 When institutions corrupt inquiry

A research system becomes less objective when it systematically weakens the routes by which errors can be found. Warning signs include:

- selective publication and undisclosed outcome switching;
- opaque data processing and unavailable code;
- conflicts of interest that affect design or interpretation;
- exclusion of relevant populations from research;
- status-based dismissal of criticism;
- incentives that reward novelty while penalizing correction.

These are not merely ethical defects. They are epistemic defects because they alter what evidence becomes visible and how it is weighted. A conflict of interest does not prove a conclusion false, but it changes the expected reliability of the process and should trigger stronger independent checks.

8.7 Objectivity as social self-correction

The view from nowhere is impossible. A more realistic ideal is a community that can locate its own partiality. RCR calls this *reflexive objectivity*. It has four components:

1. the perspective and values of inquiry are stated;
2. affected and relevantly situated people can challenge the framing;
3. independent criticism can inspect the evidence and analysis;
4. revisions are judged by their improved contact with the target, not only by their political convenience.

Reflexive objectivity is not guaranteed by diversity alone. A diverse group can share the same measurement error. Nor is objectivity guaranteed by disinterestedness alone. A narrow research agenda can be very carefully pursued. The point is to connect social inclusion with world-directed tests.

Chapter summary

Scientific knowledge is socially produced, and social organization can improve or corrupt objectivity. Division of labour, expertise, criticism, and public records create correction capacity. Merton's norms are valuable ideals but require attention to power and question selection. Longino, Harding, and Haraway show why critical pluralism and situated perspectives can reveal ignored phenomena. Values legitimately shape questions, concepts, risk thresholds, and institutions; they should not decide evidential

conclusions by fiat. Objectivity is therefore reflexive: it includes the ability of inquiry to expose its own partiality.

Objections and discussion.

1. Is a scientific community objective if it is highly diverse but has no effective correction mechanisms?
2. Which values should be public in a medical or policy study?
3. How can standpoint insights be tested without reducing them to identity claims?

Further reading. Read Longino (1990), Harding (1991), Haraway (1988), Anderson (2004), Douglas (2009), Goldman (1999), and the literature on social dimensions of scientific knowledge.

Uncertainty, Replication, and Demarcation

Chapter guide. This chapter examines how scientific systems learn from error. It compares Bayesian and error-statistical approaches, explains why replication is more than repeating a number, and develops a functional account of demarcation. The goal is not certainty, but reliable correction under uncertainty.

Fallibility is not a defect that science can eliminate. It is a condition that science can manage better or worse. A reliable inquiry does not promise that its claims are never wrong. It makes it possible to discover where they are wrong, how much depends on them, and what should replace them.

This shift matters because scientific reliability is often discussed as if one test could settle it. A p -value, a confidence interval, a Bayesian posterior, a replication, or an expert consensus can each be informative. None is self-interpreting. Objectivity lies in the relation among design, error control, evidence, and revision.

9.1 Probability as rational change

Bayesian epistemology gives a precise rule for coherent belief revision. For hypothesis H and evidence E ,

$$\mathbb{P}(H \mid E) = \frac{\mathbb{P}(E \mid H)\mathbb{P}(H)}{\mathbb{P}(E)}.$$

The rule has two virtues. It forces the investigator to compare a hypothesis with alternatives, and it distinguishes evidence from belief. The likelihood $\mathbb{P}(E \mid H)$ is not the same as the posterior probability $\mathbb{P}(H \mid E)$. A surprising result can be strong evidence for H over H_0 even when H remains improbable in an absolute sense.

The rule has limits. Priors can encode relevant background knowledge, but they can also encode unjustified assumptions. Likelihoods depend on the model of data collection and analysis. If an investigator tests many hypotheses and reports only the surprising one, the reported evidence does not have the likelihood that a preregistered test would have had.

9.2 Error-statistical learning

Deborah Mayo emphasizes learning from error: a result is informative when a procedure would probably have produced a different result if a relevant error were present (Mayo, 1996). This is a frequentist and experimental complement to Bayesian reasoning. It focuses on how a design probes error rather than on the subjective degree of belief in a hypothesis.

The two approaches need not be enemies. Bayesian updating answers how beliefs should change given a specified model. Error-statistical analysis asks whether the procedure has good capability to detect errors or discrepancies. RCR uses both as diagnostic tools:

- use likelihoods and priors to compare claims transparently;
- use error analysis to ask what the design would have detected;
- report model uncertainty and the possibility that the analysis was selected after seeing the data.

9.3 Replication is not a magic word

A replication repeats a study or result under some relation of similarity. There are several kinds:

1. **Exact replication:** repeats procedures as closely as possible.
2. **Conceptual replication:** tests the same claim using a different operationalization or design.
3. **Robustness test:** changes a condition that should not matter if the mechanism is stable.
4. **Extension:** tests whether the claim travels to a new population, scale, or context.

A failed replication can have several interpretations. The original effect may be false, the replication may have lower power, the procedures may differ in a causally relevant way, the effect may be context-sensitive, or the original estimate may have been inflated by selection. The correct response is not to declare a field true or false from one replication rate. It is to inspect the conditions under which the effect appears.

The Open Science Collaboration's 2015 project attempted replications of 100 psychological studies. It reported that replicated effects were, on average, smaller than original effects and that there was no single standard for replication success ([Open Science Collaboration, 2015](#)). The study is valuable precisely because its interpretation requires more than one headline number. It reveals how publication selection, statistical power, measurement, and context interact.

9.4 Uncertainty and the replication ladder

RCR proposes a replication ladder. A claim receives stronger support as it survives increasingly independent challenges:

1. reanalysis of the same data;
2. independent code or analytical pipeline;
3. new data with the original design;
4. different measurement or instrument;
5. different population or setting;
6. intervention on the mechanism;
7. use in a successful new prediction or application.

The ladder is not a demand that every claim reach the final rung. A historical claim may not permit intervention. A mathematical theorem has a different structure. A medical treatment may be ethically tested only within a narrow population. The ladder is a way to specify what kind of support has been achieved, not a universal ranking of sciences.

9.5 The demarcation problem

The demarcation problem asks how to distinguish science from non-science or pseudoscience. Popper proposed falsifiability. Lakatos emphasized progressive problem shifts. Laudan argued that the search for a single demarcation criterion had failed and that the more useful question is whether claims are reliable, problem-solving, and evidentially supported ([Laudan, 1977; 1983](#)).

A binary boundary is attractive for public decisions, but scientific practices are too diverse for one property to do all the work. Some mature sciences make few direct predictions; some non-scientific practices generate testable predictions. A field can be methodologically serious in one period and degenerate in another.

RCR replaces a single boundary with a profile of epistemic performance.

A claim or practice is more science-like when it:

- creates risky contact with a target;
- records methods, data, and uncertainty;
- exposes itself to independent criticism;
- updates or abandons claims when evidence requires it;
- separates empirical results from value and metaphysical interpretation;
- supports transportable predictions, interventions, or explanations.

This is a functional demarcation criterion. It does not say that science is whatever produces useful outcomes. It says that the institutional practice deserves scientific authority to the extent that it maintains correction capacity under world-sensitive constraints.

9.6 Scientism and the misuse of science

Scientism is not simply confidence in science. It is the unwarranted extension of scientific authority beyond the question, method, or evidence that supports it. A physical law does not by itself determine an ethical obligation. A statistical regularity does not settle a political priority. A model of human behaviour does not eliminate interpretation, meaning, or agency.

RCR rejects scientism for the same reason it rejects relativism: both collapse distinctions. Science is authoritative for claims that its methods can constrain. It is not a universal replacement for history, ethics, political judgment, or lived interpretation. This limitation protects science from being asked to prove what its methods cannot address, and protects other forms of reasoning from being dismissed by an illegitimate appeal to scientific status.

9.7 Reliability as a system property

A single true claim can emerge from an unreliable process by luck. A single false claim can emerge from a generally reliable process by chance. The objectivity of science is therefore not exhausted by the truth value of individual outputs. It includes the design of a system that tends to correct errors over time.

A system's reliability depends on:

1. the quality of its measurements and models;

2. the independence and competence of its critics;
3. the incentives attached to novelty and correction;
4. the visibility of negative results and failed replications;
5. the capacity to revise concepts, not only parameter estimates.

The fifth point is crucial. A field can replicate a measurement reliably while measuring the wrong construct. Objectivity sometimes requires changing the question.

Correction principle. The epistemic standing of a scientific claim depends not only on present support but on the quality of the process that could reveal its failure. A claim with moderate evidence and strong correction capacity can be more trustworthy than a claim with impressive evidence inside an opaque and uncorrectable system.

Chapter summary

Probability formalizes belief revision; error statistics asks what a procedure could have detected; replication tests stability under controlled variation; and demarcation is best understood as a profile of correction capacity rather than one binary test. Scientific authority should be proportional to the quality of world-sensitive error correction. This supports fallibilism without relativism and confidence without scientism.

Objections and discussion.

1. When does a failed replication count against a claim?
2. Is a transparent but imperfect method preferable to a highly accurate method whose assumptions cannot be inspected?
3. Can a practice be scientific in one domain and non-scientific in another?

Further reading. See Mayo (1996), the Open Science Collaboration (2015), Ioannidis (2005), Laudan (1977, 1983), and Reiss and Sprenger (2014).

The Ontology of Reflexive Constraint Realism

Chapter guide. This chapter states the ontological side of Reflexive Constraint Realism (RCR). It defines reality, entities, processes, structures, laws, levels, and social facts, then gives the framework's axioms and graded commitment ladder. The aim is a defensible realism that does not pretend that science supplies a complete inventory of being.

RCR begins with a modest but non-negotiable ontological claim: the world does not depend on human inquiry for its existence. Stars existed before astronomy. Biological reproduction preceded genetics. People and institutions can affect one another before social scientists classify them. A mind-independent world is not necessarily a world described by one final vocabulary. It is a world whose behaviour can resist, frustrate, and correct our descriptions.

The word *constraint* is central. A constraint is a stable restriction on what can happen, or a stable tendency that makes some changes more likely than others, under specified conditions. A wall constrains motion. A chemical bond constrains molecular configurations. A legal rule constrains behaviour through sanctions and expectations. A statistical relationship can express a constraint when it remains stable under relevant changes. The concept is broad enough to travel across sciences but precise enough to require a scope.

10.1 Generative reality

The ontology of RCR is *generative*: reality contains processes that produce events, not only a collection of regular sequences. A mechanism, field, institution, or ecological relation matters because it makes a difference to what can happen. Generativity does not require determinism. A stochastic

process can generate a distribution. A social institution can make behaviour more likely without forcing each individual action.

The ontology is also *layered*. Physical processes, organisms, persons, institutions, and ecosystems can be dependent on one another without being interchangeable descriptions. A higher-level entity is real when it has stable powers, boundaries, or relations that make a difference to outcomes, even if it is physically realized by lower-level processes.

Finally, the ontology is *open*. Science may discover new kinds of entity, and current categories may be replaced. The framework should therefore state what warrants commitment without claiming to know the final catalogue of reality.

10.2 Eight axioms

The following axioms are methodological-metaphysical starting points. They are not self-evident truths from which all science can be deduced. They are commitments that make the project of objective inquiry intelligible and expose where it can fail.

Axiom 1: Mind-independent generativity. There is a reality whose processes do not depend for their existence on being represented by human agents. Those processes generate resistances, regularities, possibilities, and events that inquiry can discover imperfectly.

Axiom 2: Constraint plurality. The world constrains inquiry through more than one kind of relation: material interaction, causal production, mathematical structure, temporal dependence, ecological organization, and social rule. No single layer is assumed to exhaust all scientifically relevant being.

Axiom 3: Mediated access. Scientific access to reality is always mediated by concepts, instruments, models, language, and institutions. Mediation does not imply invention; it creates a route through which independent constraints can act on claims.

Axiom 4: Fallibility. Every finite scientific claim is revisable. Confidence can be high without being certain, and ontological commitment

can be strong without being final.

Axiom 5: Invariance matters. A claim gains epistemic and ontological support when its relevant consequences remain stable under admissible changes of observer, instrument, representation, intervention, or context.

Axiom 6: Situated inquiry. Every inquiry begins from a location in a social, material, and historical system. The location affects what becomes visible, but it need not determine what is true.

Axiom 7: Public corrigibility. Objective inquiry requires practices through which competent others can inspect, challenge, reproduce, and revise a claim. Private certainty is not a substitute for public correction.

Axiom 8: Scope-bound truth. A claim is true only with respect to the objects, variables, scales, interventions, and background conditions that its evidence supports. Generality must be earned by transport.

Axioms 1 and 3 must be held together. If one accepts only Axiom 1, one risks naive realism: the world is real, so observations are treated as transparent. If one accepts only Axiom 3, one risks constructivist relativism: all access is mediated, so reality becomes indistinguishable from representation. Axioms 5 through 8 explain how mediation can be disciplined.

10.3 What scientific entities are

An entity is not simply whatever appears as a noun in a theory. RCR uses three tests.

1. **Boundary test:** Can the entity be distinguished from neighbouring processes by a stable criterion?
2. **Difference test:** Does changing or removing the entity make a reliable difference to outcomes?
3. **Convergence test:** Do independent measurements, models, or

interventions support the entity's role?

An entity can pass these tests without being fundamental. A gene may be a molecular sequence, a regulatory function, or a population-level unit depending on the question. A firm is not a biological organism, but it can be a real social entity if it has legally recognized boundaries, resources, decision procedures, and causal effects.

RCR therefore distinguishes *realization* from *reduction*. To say that a higher-level entity is realized by lower-level processes is to say that the higher-level entity depends on them. To say that it is reducible is to say that its explanatory and causal role can be replaced without loss. The first does not entail the second.

10.4 Structures and relations

A structure consists of elements together with relations and transformations. Scientific theories often support relations more strongly than they support the intrinsic nature of elements. A model may establish that changes in X alter Y with a certain pattern without determining whether X is a fundamental property or a useful aggregate.

The realist commitment to structure must be disciplined by relevance. A mathematical isomorphism alone does not show that two physical systems share the same ontology. The relation must survive a mapping between model and target, and it must matter to prediction, intervention, or explanation.

10.5 Laws as modal constraints

RCR treats laws of nature as candidates for stable modal constraints: they describe what would happen under a range of admissible changes, not merely what has happened so far. A law-like claim can be probabilistic, approximate, or scale-specific. It becomes law-like when its counterfactual and interventional consequences are stable enough to guide inquiry.

This view avoids the choice between universal regularity and mere description. A law can be a real feature of a domain without being the only law that exists or without applying outside the regime in which its variables are well-defined.

10.6 The ontology of social facts

Social facts are not less real because they depend on human activity. Money, property, citizenship, corporations, and scientific institutions have histories

and rules, but they constrain behaviour and distribute resources. Their existence is relational and reflexive: they persist because people recognize, enforce, or enact them, yet once established they are not freely adjustable by each individual.

The ontology of social science must therefore include at least:

- material bodies and environments;
- agents with capacities and interpretations;
- rules, norms, and institutions;
- networks of dependence and power;
- historical trajectories that alter future possibilities.

A social explanation should not collapse a norm into a molecule, but neither should it detach the norm from the people, artifacts, and sanctions that realize it.

10.7 Graded ontological commitment

RCR recommends a commitment ladder. The ladder is not a probability scale. It records what sort of support has been earned.

The ladder protects against two errors. It prevents a useful model from being treated as a literal inventory. It also prevents anti-realism from ignoring the epistemic significance of interventions that could not succeed if the relevant entity or structure were absent.

10.8 A layered map of reality




Figure 10.1: A non-reductive layered ontology. Higher layers are realized by lower processes while feeding back through organization, action, and selection.

Level	Commitment	Typical warrant
L0: useful posit	Use the term without belief in its independent reality	Predictive or computational utility
L1: empirical pattern	Believe that a stable phenomenon occurs	Convergent measurement and replication
L2: structural commitment	Believe that a relation or invariance is real in a domain	Cross-model stability and limiting relations
L3: causal commitment	Believe that an entity, process, or mechanism makes a difference	Intervention, counterfactual robustness, mechanism
L4: domain ontology	Treat a family of entities and relations as real in a mature domain	Long-run convergence, transport, integration, and correction

Table 10.1: The RCR commitment ladder. Movement upward requires stronger and more independent constraints; movement downward remains possible after counterevidence.

The diagram is not a ladder of value or a claim that reality has exactly four levels. It is a reminder that scientific explanations can move across levels. A social institution is materially realized, psychologically interpreted, and historically reproduced. A disease is biological, embodied, measured, and institutionally classified. Objectivity depends on preserving the causal links between levels rather than pretending that one level can answer every question.

Chapter summary

RCR posits a mind-independent, generative, layered, and open reality. Its basic unit is the constraint: a stable restriction or tendency that survives specified changes. Entities earn commitment through boundary, difference, and convergence tests. Structures are real when relevant relations remain stable across representations and interventions. Laws are modal, scope-bound constraints. Social facts are real relational entities. The eight axioms and the commitment ladder make realism graded, fallible, and accountable.

Objections and discussion.

1. Can an entity be real without being fundamental?
2. What would count as evidence that a higher-level social mechanism has causal powers?
3. Does the commitment ladder describe degrees of belief, degrees of truth, or degrees of warrant?

Further reading. Compare the layered approach here with Chakravartty (2007), Cartwright (1999), Craver (2007), Batterman (2002), Oppenheim and Putnam (1958), and the literature on reduction and emergence.

The Epistemology of Reflexive Constraint Realism

Chapter guide. This chapter gives RCR an epistemology. It defines evidence, states inferential rules, explains how belief should change, and proposes criteria for comparing theories. The framework does not replace Bayesian, causal, experimental, or interpretive methods; it specifies how their results can be assembled into an objective but fallible claim.

The ontological claim that reality constrains inquiry is not enough. A world can be real while our beliefs about it are systematically wrong. RCR's epistemology therefore asks how constraints become reasons. A reason is not simply an event that occurs before belief. It is evidence when the event discriminates among claims under a defensible measurement and inferential model.

11.1 Claims, targets, and alternatives

Every scientific claim should be represented with five fields:

1. **Target:** the entity, process, relation, or population at issue;
2. **content:** what is being asserted about the target;
3. **scope:** the time, place, scale, population, and intervention conditions;
4. **alternatives:** rival explanations or models;
5. **warrant:** the evidence and reasoning supporting the claim.

A sentence such as “ X causes Y ” is incomplete. It does not say which X , which Y , in what population, by what intervention, or against which

confounders. A scoped claim can be tested. A vague claim can survive by changing its meaning.

11.2 Evidence as constraint

Evidence E supports a claim H when E is more expected under H than under relevant alternatives, or when E forces a revision of those alternatives. The Bayesian expression is the likelihood ratio:

$$\text{LR}(E; H_1, H_0) = \frac{\mathbb{P}(E \mid H_1)}{\mathbb{P}(E \mid H_0)}.$$

A likelihood ratio greater than one favours H_1 over H_0 ; a ratio less than one favours H_0 . The ratio does not establish either hypothesis as true. It measures the evidential direction and strength relative to a model.

In practice, $\mathbb{P}(E \mid H)$ includes assumptions about sampling, measurement, missingness, and analysis. RCR therefore requires that the evidence statement carry its measurement path and scope conditions with it.

Evidence rule. Do not ask whether a result is evidence in isolation. Ask: evidence for which claim, against which alternative, under which measurement path, and with what expected error if the claim were false?

11.3 The inferential rules

The following rules are the operational core of RCR. They are norms for reasoning, not algorithms that mechanically settle every case.

- R1. Scope before strength.** State the claim's target and scope before assigning confidence. A highly precise claim with an undefined target has little epistemic content.
- R2. Path traceability.** Trace the path from target to datum to conclusion. Separate observation, transformation, model assumption, and interpretation.
- R3. Alternative exposure.** Compare a claim with serious alternatives. A result that fits H but also fits many rivals has limited discriminating force.
- R4. Risk preference.** Prefer tests that create outcomes likely to differ across hypotheses. Confirmation is stronger when the test was not designed to guarantee success.

- R5. Invariance search.** Vary observers, instruments, operational definitions, samples, models, or interventions where the claim says the result should remain stable.
- R6. Causal separation.** Do not infer causal adequacy from predictive fit alone. Require design, intervention, mechanism, or a clearly defended counterfactual model.
- R7. Transport discipline.** Do not move a result to a new setting without identifying shared mechanisms and checking changed distributions.
- R8. Commitment proportionality.** Match ontological commitment to the strongest layer of evidence actually established.
- R9. Revision readiness.** State what evidence would lower confidence or change the model. A claim insulated from revision is not strengthened by confidence.
- R10. Social accountability.** Treat competent, relevant criticism as a source of information about hidden assumptions, not merely as a political obstacle.

The rules can conflict. A high-risk test may be unethical. An exact replication may be impossible. A new operationalization may alter the target. In such cases, the conflict should be reported rather than hidden. RCR is a framework for making trade-offs visible, not a claim that every trade-off has one correct answer.

11.4 Belief revision

Belief revision should be continuous, comparative, and calibrated. Continuous does not mean that all beliefs must be represented by exact numbers. It means that the language of confidence should reflect degrees of support. Comparative means that confidence in one claim depends on how it performs against alternatives. Calibrated means that predictions are checked against outcomes.

A useful reporting format separates:

- prior background assumptions;
- new evidence;
- the likelihood or expected error under each alternative;
- the updated confidence;
- remaining uncertainty and scope limitations.

This format is compatible with Bayesian updating, confidence intervals,

likelihoodist reasoning, error statistics, or qualitative comparative inference. The common requirement is that the route from evidence to belief be auditable.

11.5 Theory comparison

Theories should not be ranked by one virtue alone. A simple theory can be too simple. A detailed theory can overfit. A highly predictive theory can lack a causal mechanism. A metaphysically rich theory can make no risky difference.

RCR proposes a comparative profile with seven dimensions:

1. empirical fit to established phenomena;
2. novel or risky success;
3. structural and causal coherence;
4. robustness under independent changes;
5. explanatory reach and integration;
6. simplicity relative to scope and complexity;
7. transparency, corrigibility, and transportability.

For an informal representation, let a theory T have profile

$$\mathbf{V}(T) = (F, N, C, R, X, S, O),$$

where the letters stand for fit, novelty, coherence, robustness, explanatory reach, simplicity, and objectivity of process. The profile is not a single score. In some decisions one dimension can dominate: a medical treatment may require safety even if its explanatory theory is incomplete. In others, structural coherence or novel prediction may matter more.

A weighted score can be used when a purpose requires one:

$$Q(T) = \sum_{j=1}^7 w_j V_j(T), \quad w_j \geq 0, \quad \sum_{j=1}^7 w_j = 1.$$

Here w_j are explicit weights chosen for a stated purpose. The equation does not create objectivity by arithmetic. It prevents hidden trade-offs from masquerading as neutral judgment.

Dimension	Strong performance	Typical warning sign
Fit	Accounts for reliable phenomena within error	Selective use of data
Novelty	Predicts or reveals outcomes not used to tune the theory	Post hoc accommodation
Coherence	Links mechanisms, equations, and concepts without contradiction	Ad hoc patches
Robustness	Survives instruments, samples, and model changes	Fragility to one pipeline
Reach	Connects domains or scales with clear bridges	Vague unification
Simplicity	Achieves scope without needless complexity	Underfitting or oversimplification
Process objectivity	Makes criticism and revision possible	Opaque methods or status protection

Table 11.1: A comparative profile for theory choice. No dimension is sufficient on its own.

11.6 Explanation and evidence are distinct

A theory may explain a phenomenon but fit data poorly because the measurements are noisy. A statistical model may fit and predict accurately without identifying the mechanism. A mechanism may be true while a quantitative parameter is wrong. RCR avoids forcing these into one score by reporting separate warrants.

This distinction is also useful for scientific communication. A paper should not describe an associational analysis as a causal explanation merely because the proposed story is plausible. It should not describe a computational possibility as an empirical prediction until the bridge to the target is identified. It should not describe a normatively preferred outcome as a scientific conclusion without separating values from evidence.

11.7 Objectivity criteria

A claim qualifies as objective in the RCR sense when it satisfies the following criteria to a degree appropriate to its field:

1. **World resistance:** the claim confronts effects that investigators

cannot freely choose.

2. **Cross-perspective convergence:** relevantly different observers or methods identify the same pattern.
3. **Measurement accountability:** the path from target to result is traceable and error-aware.
4. **Intervention or counterfactual grip:** the claim supports a well-defined change relation where one is possible.
5. **Model transparency:** idealizations and bridge assumptions are stated.
6. **Critical independence:** competent others can inspect and challenge the claim.
7. **Scope honesty:** the claim does not travel further than its evidence.
8. **Revision capacity:** the inquiry can detect and incorporate error.

These criteria are not a checklist that produces certainty. They are dimensions along which a research system can improve. A claim may score strongly on measurement and weakly on intervention because the target is historical. It may score strongly on causal grip and weakly on transport because the setting is unique.

11.8 What makes the epistemology reflexive

Reflexivity means more than adding a disclaimer about bias. It requires studying how the inquiry's own categories and procedures change the target, the evidence, or the community. In social science, a classification can become an incentive. In medicine, a diagnostic label can affect treatment. In physics, a detector design defines which events are accessible. In all cases, the inquiry should ask what it has made visible and what it has excluded.

A reflexive epistemology is not self-defeating. The fact that an inquiry has conditions does not imply that its conclusions are only about those conditions. It means that a conclusion earns broader authority when it survives changes to the conditions that are not supposed to matter.

Chapter summary

RCR treats evidence as a claim-relative constraint, not as a free-floating fact. Its ten inferential rules require scope, traceability, alternatives, risk, invariance, causal separation, transport discipline, proportional commitment, revision readiness, and social accountability. Theory comparison uses a profile of fit, novelty, coherence, robustness, reach, simplicity, and

process objectivity. The criteria of objectivity are graded, field-sensitive, and designed to make error discoverable.

Objections and discussion.

1. Should a theory with stronger fit but weaker causal explanation be preferred for prediction?
2. Which of the ten rules is most difficult to satisfy in social science?
3. Can explicit values improve theory comparison rather than bias it?

Further reading. Compare the updating logic here with Bayes (1763), Reichenbach (1938), Mayo (1996), and the causal frameworks of Pearl (2009) and Woodward (2003).

Scientific Progress and Theory Choice

Chapter guide. This chapter asks in what sense science progresses if theories are revised and models are idealized. It distinguishes accumulation from improvement, considers underdetermination and unconceived alternatives, and explains how RCR treats theory change as the enlargement and stabilization of constraint-sensitive inference.

If objectivity is fallible, what makes one theory better than another? A later theory may be less simple, less intuitive, or less familiar. It may reject entities once considered real. It may also predict more accurately, explain more phenomena, support new interventions, and clarify its own limitations. Scientific progress is not the absence of change. It is change that increases the reliability, range, precision, or depth of inquiry.

12.1 Against simple accumulation

The accumulation picture says that science adds facts to a fixed stock of knowledge. Kuhn showed why this is inadequate: revolutions change concepts, problems, and standards (Kuhn, 1962). But the opposite picture, in which each theory is a self-contained world, is also inadequate. Later science often preserves measurements, limiting cases, mathematical relations, instruments, and practical controls.

RCR identifies five forms of progress:

1. **Empirical progress:** more reliable phenomena are established and fewer important results remain unexplained.
2. **Predictive progress:** the theory yields new, risky, and calibrated predictions.
3. **Causal progress:** the theory identifies more stable mechanisms and intervention boundaries.

4. **Representational progress:** models preserve relevant structure across wider ranges or provide better approximations.
5. **Reflexive progress:** the inquiry becomes better at identifying its own errors, assumptions, and social exclusions.

These dimensions can conflict. A more accurate model may be less interpretable. A broader theory may weaken local fit. A new measure may reveal that an old effect was an artifact. Progress is therefore multidimensional and must be reported as a profile.

12.2 The problem of underdetermination

For any finite set of observations, many models may fit. This is underdetermination. It follows from the fact that data constrain hypotheses but do not usually select one uniquely. Duhem's holism and Quine's web of belief make the problem vivid (Duhem, 1906; Quine, 1951).

The existence of alternatives does not imply that theory choice is arbitrary. Researchers can generate new constraints by measuring a different variable, designing a discriminating experiment, comparing interventions, or testing novel domains. The appropriate response to underdetermination is not to declare victory for the preferred theory. It is to ask what further test would separate the live alternatives and whether the difference matters to the question.

Kyle Stanford's discussion of unconceived alternatives strengthens the warning. Even when scientists compare the theories they have imagined, they may fail to consider a future theory that preserves present successes while reinterpreting their ontology (Stanford, 2006). RCR accepts that no current commitment is immune to future alternatives. It responds through graded commitment and a distinction between secure constraints and provisional interpretations.

12.3 Convergence and robustness

Convergence is not mere agreement. Two studies that use the same flawed instrument and code may be correlated failures. Convergence is epistemically strong when the routes share the target but differ in potential errors.

Examples of productive convergence include:

- a physical quantity measured by distinct instrument principles;
- a biological mechanism supported by molecular, cellular, and organ-

ismic interventions;

- a climate attribution supported by observations, physical models, energy-balance reasoning, and independent indicators;
- a policy effect supported by randomization, quasi-experiment, mechanism, and implementation evidence.

Robustness tests should be chosen to vary what is not supposed to matter and to preserve what is supposed to matter. If a claim survives only when every analytic choice is held fixed, its objectivity is weaker than its apparent precision suggests.

12.4 Simplicity and complexity

Simplicity is often treated as a reason to prefer one theory. But simple in what sense? A theory can have fewer parameters but a more complicated ontology. A concise model can hide a complex data-processing pipeline. A complex model can be more honest if the target is heterogeneous.

RCR treats simplicity as a resource for transport and criticism, not as a metaphysical law. A model is preferable when its complexity is proportionate to the structure it must represent and when the added complexity earns new constraint-sensitive success. This is close to the methodological idea behind penalizing overfit models, but it is broader: complexity includes hidden assumptions, institutional dependencies, and unreported flexibility.

12.5 The role of explanation

Scientific progress is not only increased prediction. A black-box system can forecast accurately but fail to reveal which intervention would change the outcome. An explanation makes a difference to what questions can be asked next.

Philip Kitcher connects scientific progress with the organization and unification of knowledge (Kitcher, 1993). A good explanation can reduce dependence on isolated facts and improve the ability to derive new consequences. But unification can be forced. RCR counts explanatory reach only when the bridges between domains are explicit and survive independent tests.

12.6 Theory change and retained constraints

A theory change is progressive when it replaces a representation while retaining or improving constraints that mattered. Consider a sequence of models $M_1, M_2, \dots, M_k$. A progress relation is not simply M_{j+1} has more

equations than M_j . It can include:

$$\text{Prog}(M_{j+1}, M_j) = \text{Retain}(R) + \text{Extend}(D) + \text{Correct}(A),$$

where R is a set of retained relations, D is newly constrained domain behaviour, and A is a set of known anomalies or assumptions corrected by the new model. The expression is conceptual, not a numerical theorem.

A later theory can therefore be better even if it says that an older entity was not fundamental. It may retain the old theory as a limiting approximation, explain why the approximation worked, and predict where it fails. This is one meaning of objective progress.

12.7 Research programmes and the future

A scientific field should not aim only to confirm its current framework. It should cultivate a portfolio of activities:

1. normal work that improves precision and solves established puzzles;
2. anomaly investigation that tests the framework's boundaries;
3. alternative model construction that exposes hidden assumptions;
4. intervention and replication that probe causal stability;
5. conceptual analysis that revises categories when measurements conflict;
6. institutional reform that improves correction capacity.

A field that lacks anomaly work can become dogmatic. A field that lacks normal work cannot stabilize knowledge. A field that lacks alternatives cannot know whether its theory is overfitted. A field that lacks institutional reform can identify errors without being able to correct them.

Progress principle. Scientific progress is the cumulative enlargement and stabilization of world-sensitive constraints, accompanied by improved ability to identify scope, error, and alternative representations. It need not preserve a fixed ontology or produce a final theory.

Chapter summary

Progress is multidimensional: empirical, predictive, causal, representational, and reflexive. Underdetermination is real but can be reduced by discriminating tests and independent routes of evidence. Convergence

requires different potential error structures, not repeated agreement from one pipeline. Simplicity is a methodological resource rather than a metaphysical guarantee. Theory change is progressive when it retains important constraints, extends their domain, corrects anomalies, and improves the ability to find errors.

Objections and discussion.

1. Can a theory be progressive if it increases prediction but decreases interpretability?
2. How can researchers search for alternatives they have not yet imagined?
3. Which forms of progress matter most in policy-relevant science?

Further reading. See Lakatos (1978), Laudan (1977), Kitcher (1993), Stanford (2006), Cartwright (1999), and the literature on structural realism and scientific progress.

Cases: Testing Reflexive Constraint Realism

Chapter guide. This chapter tests RCR across different sciences. Each case asks what was measured, what constraints were established, which parts of the successful representation were idealized, and what the case does not justify. The cases include physics, biology, medicine, psychology, climate science, economics, and social science.

A philosophy of science should be answerable to science as it is practiced. The following cases are not offered as decisive proof of RCR. They are stress tests. Each case has a different pattern of access, evidence, intervention, and institutional organization. If one framework can illuminate all of them without pretending that their methods are identical, it has earned some credibility.

13.1 Physics: atoms, electrons, and theory change

Nineteenth-century debates about atoms show why ontological commitment should be graded. Kinetic and chemical theories used atoms to explain regularities, but some scientists treated atoms as calculational devices rather than established entities. Einstein's 1905 analysis of Brownian motion connected observable particle movement with molecular-kinetic theory. Jean Perrin later used several quantitative investigations of Brownian motion and related phenomena to support the atomic theory ([Einstein, 1905](#); [Perrin, 1913](#)).

The philosophical significance is not that one paper proved atoms in a deductive sense. It is that independent relations converged: diffusion, sedimentation, viscosity, and particle motion could be connected by a common molecular framework. The framework generated quantitative constraints that were not simply restatements of a single observation. Under

RCR, atoms move from useful posit toward strong causal and structural commitment because the same posit supports multiple measurement paths.

The electron provides a related case. J. J. Thomson's 1897 work on cathode rays interpreted their deflection as evidence for negatively charged constituents smaller than atoms. Subsequent measurements, including Millikan's oil-drop work, connected charge and mass through distinct experimental arrangements. The history is not a story of one experiment silently revealing an entity. It is a sequence of instruments, calibrations, controversies, and theoretical interpretations in which a stable causal role emerged.

The Michelson–Morley experiment of 1887 illustrates a different point. It did not by itself refute every version of an ether theory or uniquely establish special relativity. Its result became significant within a larger network of optics, mechanics, and later theoretical developments (Michelson and Morley, 1887). This is a case where the Duhem–Quine problem is not an excuse for ignoring evidence. The experiment constrained a family of models and contributed to a reorganization of theoretical structure.

RCR verdict. Physics supports realism about certain entities and relations when they survive independent measurement and intervention, but it does not support the assumption that every historical interpretation of those entities was correct. What is retained may be a relation, a field equation, a limiting case, or an interventionally stable capacity rather than a complete ontology.

13.2 Biology: the structure of DNA

The 1953 proposal by James Watson and Francis Crick represented DNA as a double helix with paired bases (Watson and Crick, 1953). The result was not deduced from a single image. It combined chemical constraints, model-building, and X-ray diffraction evidence, including work by Rosalind Franklin and Raymond Gosling. The published structure was a selective representation: it captured a geometry and pairing relation without being a complete theory of replication, gene regulation, or development.

This case displays all four layers of scientific success.

1. **Correspondence:** DNA has a molecular structure that constrains its chemical behaviour.
2. **Structural accuracy:** the double-helical arrangement preserves relations among the backbone and base pairs.

3. **Interventional adequacy:** later molecular biology could alter, sequence, and manipulate nucleic-acid systems in ways that depended on the structure.
4. **Pragmatic adequacy:** the representation supported new research questions and technologies.

The case also shows why social epistemology matters. The historical credit attached to a discovery does not perfectly track the distribution of data, labour, and interpretation. A theory can be scientifically strong while the institutional history of recognition is unjust. Epistemic evaluation should therefore separate the truth of a structural claim from the fairness of the community that produced and rewarded it.

13.3 Medicine: the MRC streptomycin trial

The Medical Research Council's 1948 controlled investigation of streptomycin for pulmonary tuberculosis is an important case of experimental and institutional design ([Medical Research Council, 1948](#)). The trial used allocation procedures intended to prevent foreknowledge of treatment assignment and compared outcomes under a defined protocol. The result did not establish that streptomycin was a complete solution to tuberculosis. Resistance, toxicity, disease severity, and combination therapy remained relevant.

Its philosophical importance lies in the separation of causal question from clinical enthusiasm. The question was not merely whether patients receiving streptomycin appeared to improve. It was whether outcomes differed under a controlled comparison that reduced selection bias. Randomization did not make the study theory-free. It required definitions of disease, treatment, outcome, follow-up, and missingness. But it created a constraint that a persuasive narrative alone could not easily evade.

Under RCR, the trial strongly supports a scoped causal claim: under the studied conditions, the treatment affected specified outcomes in the studied population. It does not license an unqualified claim that the drug works in all settings, for all forms of disease, or as a universal mechanism. The intervention boundary matters.

13.4 Psychology: replication and context

The Open Science Collaboration's 2015 project attempted to replicate 100 psychological studies and reported that replicated effects were, on average, smaller than the original effects ([Open Science Collaboration, 2015](#)). The project is often summarized as a simple failure rate. That summary is

inadequate. Replication success can be evaluated by significance, effect size, interval overlap, prediction intervals, or theoretical fit. Each criterion answers a different question.

The case illustrates three RCR principles.

First, a result is not fully objective because it passed one statistical threshold. Publication selection, flexible analysis, and measurement can inflate apparent evidence.

Second, a failed replication is evidence about the claim and the conditions, not a metaphysical verdict on a whole discipline. If an effect depends on population, context, stimulus, or implementation, a conceptual replication may reveal a boundary rather than a simple falsehood.

Third, social organization is part of epistemology. A field that rewards novelty and treats correction as failure weakens its own correction capacity. Preregistration, open materials, independent analyses, and better incentives can improve objectivity without guaranteeing uniform results.

13.5 Climate science: convergence across models and evidence

Climate science is a large-scale case of model-mediated, multi-source inference. No single thermometer, simulation, or statistical method establishes the whole climate system. The evidence includes instrumental records, paleoclimate proxies, physical theory, radiative measurements, ocean heat content, cryosphere changes, attribution studies, and model projections.

The IPCC's Sixth Assessment Synthesis Report, released in 2023, assesses that human influence has warmed the atmosphere, ocean, and land, and it expresses the assessment with calibrated language about confidence and likelihood (IPCC, 2023). The report does not remove uncertainty. It distinguishes established observed changes, attribution, projected risks, and the conditions under which projections vary.

Climate science therefore tests RCR in a domain where direct intervention on the global system is unavailable and the target is dynamic. Its objectivity comes from convergence among independent constraints and from the physical mechanisms linking greenhouse-gas concentrations, energy balance, circulation, and observed changes. Models are not literal replicas of the atmosphere. They are idealized tools whose reliability depends on hindcasts, process validation, inter-model comparison, and observation.

The case also shows why values enter without deciding the physical conclusion. Choices about acceptable risk and vulnerable populations affect

policy interpretation, but they do not determine whether human forcing has altered the climate system.

13.6 Economics: policy models and the Lucas critique

Economic models demonstrate the importance of reflexivity. A policy can change the expectations and behaviour that generated an earlier statistical relationship. Robert Lucas's critique argued that policy evaluation based on historical correlations can fail when policy changes the decision rules of agents (Lucas, 1976).

The philosophical lesson is not that economic models are unscientific. It is that transportability requires a mechanism. A model estimated under one policy regime may not predict under another if households, firms, or institutions adapt. Causal identification must therefore address expectations, equilibrium, institutional rules, and the time path of adjustment.

RCR recommends separating several claims:

- a historical association among variables;
- a structural model of decision rules;
- a causal effect under a specified intervention;
- a forecast conditional on agents and institutions not changing in unmodelled ways.

An economic model can be useful and structurally informative without literally describing every agent. Its objectivity depends on the stability of the mechanism under the policy intervention being considered.

13.7 Social science: field experiments and discrimination

Field experiments can create causal constraints in settings where laboratory control is limited. Bertrand and Mullainathan's audit study sent matched resumes with different names to employers and examined differences in callbacks (Bertrand and Mullainathan, 2004). The design was intended to isolate the effect of perceived racial identity on employer response while keeping other resume information comparable.

The study does not identify every mechanism of discrimination. Names may signal more than one social perception; employers may differ in interpretation; and callback differences do not exhaust later hiring processes.

It does, however, show how a social category can be investigated through an intervention that changes a signal while holding much of the application constant.

The case supports a scoped claim about employer responses under the study's conditions. It also demonstrates the reflexive nature of social research: categories such as race are socially produced, but their effects on decisions can be objectively studied. The reality of a social mechanism is not reduced to a biological essence or dissolved into language.

13.8 Cross-case comparison

Field	Main access	Strongest constraint	Main limit
Physics	Instrumental interaction and theory	Convergent measurement and intervention	Theory change can alter ontology
Biology	Structure, experiment, mechanism	Cross-scale causal integration	Complex context and multiple levels
Medicine	Controlled intervention	Randomized treatment comparison	Scope, ethics, heterogeneity
Psychology	Experiment, measurement, replication	Independent reproduction and design transparency	Context and measurement variation
Climate	Observation, physical model, attribution	Convergence across sources and scales	No global intervention; projection uncertainty
Economics	Historical data, models, policy variation	Mechanism under changing expectations	Reflexivity and transportability
Social science	Field experiment, records, interpretation	Scoped intervention and plural critique	Categories and institutions can change

Table 13.1: Different sciences realize objectivity through different constraint profiles.

The cases support a modest generalization. Objective science is not defined by one preferred method. It is defined by the capacity of a method and institution to create relevant constraints, expose error, and state what the result does not establish.

Chapter summary

Physics shows how entities and structures earn commitment across independent measurements. Biology shows how model-building and social credit can diverge. Medicine shows the causal value of controlled comparison and the importance of scope. Psychology shows why replication is a graded and contextual test. Climate science shows how convergence can support strong conclusions without eliminating uncertainty. Economics shows that policy changes can alter the mechanisms being estimated. Social science shows that socially constructed categories can have objectively testable causal effects.

Objections and discussion.

1. Which case provides the strongest evidence for entity realism?
2. Which case most clearly separates pragmatic usefulness from causal adequacy?
3. How should a climate or economic model report uncertainty that comes from future institutional change?

Further reading. Read the primary studies cited in the cases. For broader comparisons, consult Hacking (1983), Woodward (2003), Cartwright (1999), Craver (2007), and Kitcher (2001).

Objections, Revisions, and Limits

Chapter guide. A framework earns credibility by exposing itself to its strongest objections. This chapter asks whether RCR is merely a disguised realism, an unstable mixture of methods, a version of pragmatism, an impossible demand on institutions, or an unhelpful abstraction. Each objection produces a qualification or revision.

No account of scientific objectivity can avoid trade-offs. A framework that solves every problem by adding another virtue becomes a catalogue rather than a theory. RCR must therefore face objections that target its central ideas: constraint, convergence, intervention, reflexivity, graded realism, and public correction.

14.1 Objection 1: the framework merely renames realism

A critic may say that RCR is scientific realism with extra vocabulary. Realists already appeal to successful reference, structure, intervention, and causal power. Why introduce “constraint”?

The objection is partly correct. RCR does not claim to have invented the world, evidence, or scientific realism. Its originality lies in the architecture: it links a generative ontology to a claim-relative account of evidence, a graded commitment ladder, intervention boundaries, and social correction criteria. The term constraint is useful because it can describe what is shared by an instrument reading, a causal mechanism, a mathematical invariance, and a social rule without pretending that they are the same kind of thing.

Revision. RCR should not present itself as a wholly unprecedented metaphysics. It is an original framework and argument built from, and

sometimes opposed to, established ideas.

14.2 Objection 2: constraints are too broad to be explanatory

If everything that restricts possibilities is called a constraint, the concept may explain nothing. A wall, a law, a probability distribution, and a norm are not interchangeable.

This objection requires a taxonomy. RCR distinguishes:

- **material constraints:** physical dispositions and interactions;
- **formal constraints:** mathematical or logical relations;
- **causal constraints:** stable change-relations under interventions;
- **organizational constraints:** relations among parts in a system;
- **institutional constraints:** rules sustained by recognition, sanction, and material arrangements.

The shared concept is functional, not categorical. It says that objectivity depends on resistance to arbitrary reinterpretation. The explanation must then identify the kind of constraint and the process that sustains it.

Revision. A scientific claim should never report only that it found a constraint. It should specify the constraint type, scope, realization, and failure conditions.

14.3 Objection 3: convergence can be correlated failure

Several methods can agree because they share a hidden assumption. Different climate models may share parameterizations. Different surveys may use the same biased sampling frame. Replications may reuse the same flawed operational definition. Convergence therefore does not guarantee truth.

This is a serious objection. RCR never treats the number of agreeing studies as a sufficient measure of independence. The relevant question is whether the routes have different potential errors. Independence is epistemic and counterfactual, not merely organizational.

Revision. Convergence should be evaluated by an error-correlation audit. Researchers should state which assumptions and failure modes are shared across methods. A new method earns more weight when it changes the relevant error structure.

14.4 Objection 4: intervention is theory-laden

Hacking and interventionist philosophers may seem to promise a direct route to reality, but interventions are designed through theory. A manipulation may change several variables. A detector may be interpreted through a model. A successful intervention can therefore support multiple ontologies.

The objection is right that intervention is not a magical contact with uninterpreted reality. Its epistemic value comes from the way it changes possibilities under controlled conditions. Intervention supports a claim when the manipulation is specified, the path from manipulation to outcome is measured, and rival mechanisms predict different results.

Revision. RCR replaces “intervention proves entity” with the weaker rule that intervention raises commitment when the effect is stable under apparatus changes, not used to define the outcome, and integrated with independent evidence.

14.5 Objection 5: Bayesian updating makes objectivity subjective

Different priors can yield different posteriors. If rationality depends on prior belief, perhaps scientific objectivity is just agreement among people with similar backgrounds.

Bayesianism does not imply that every prior is equally good. Priors can encode relevant background knowledge, symmetry, domain knowledge, or institutional assumptions. They can also be criticized. Strong likelihood ratios can reduce initial differences, but not always. In cases of weak data, disagreement may remain rational.

RCR responds by separating three questions: coherence of updating, quality of background assumptions, and strength of the data’s discrimination. Public objectivity requires making priors and model choices visible, comparing sensitivity to alternative priors, and using new experiments to create stronger likelihood differences.

Revision. When priors materially affect a conclusion, a study should report a prior-sensitivity analysis rather than one apparently neutral answer.

14.6 Objection 6: social criticism cannot reach mind-independent truth

A social epistemologist may argue that standards of criticism are themselves historical and political. Why should institutional agreement tell us anything about a world independent of institutions?

The answer is that social correction is a causal part of evidence production, not a source of truth by itself. Institutions can improve or weaken access. A community's conclusion becomes more objective when its procedures expose claims to independent material constraints and when critics can identify errors. Social consensus without world resistance is not enough.

Revision. RCR must keep two levels distinct: the social conditions of inquiry and the target that constrains it. Neither level can replace the other. A just institution is not automatically a truthful one; a technically successful result does not excuse an unjust institution.

14.7 Objection 7: standpoint theory becomes identity essentialism

If some standpoints have epistemic advantages, critics may fear that membership in a group becomes a substitute for argument. This would reverse the error of exclusion without solving it.

RCR accepts the objection. Standpoint advantage is conditional. It arises when a social position exposes relations hidden from dominant institutions and when the insight is developed through critical reflection and collective testing. No identity guarantees a correct conclusion, and no identity disqualifies a person from understanding.

Revision. Standpoint claims should generate research questions and tests. They should not be treated as conclusions immune to evidence.

14.8 Objection 8: values contaminate objectivity

A strict value-free ideal might say that values should be excluded entirely. If values enter question selection, concepts, and risk thresholds, perhaps no scientific conclusion is objective.

The response is a distinction. Values can affect which questions receive resources and how uncertainty is used in decisions. They should not make a physical observation or statistical comparison true merely because it

serves an interest. Indeed, hiding values can be more corrupting than acknowledging them because it prevents criticism of the choice.

Revision. RCR requires value transparency and keeps empirical, interpretive, and normative claims in separate statements. It does not demand that all scientific work be politically neutral in its consequences.

14.9 Objection 9: the commitment ladder is arbitrary

Why should a posit move from a useful tool to a structural or causal commitment after a certain kind of evidence? The ladder might merely encode the author's preferences.

The ladder is not a universal threshold. It is a vocabulary for reporting what kind of support has been established. A community may disagree about whether a result meets the intervention standard. That disagreement is precisely the kind of dispute the framework wants to expose.

Revision. The ladder should be used comparatively and revisably. Reports should state which rung is claimed, what evidence supports it, and what evidence would move the claim down.

14.10 Objection 10: the framework is too demanding for real research

Few studies satisfy every criterion. Scientists have limited time, budgets, ethical constraints, and access to data. If RCR demands independent instruments, interventions, replication, and social inclusion everywhere, it will paralyze inquiry.

This objection is decisive against a rigid checklist. RCR is instead a calibration framework. The required evidence depends on the claim and its stakes. A low-stakes exploratory model can have modest correction capacity. A high-stakes medical or policy claim warrants stronger independent tests. A historical claim needs source criticism and causal trace analysis rather than randomization.

Revision. Objectivity requirements should be proportional to the claim's risk, generality, and irreversibility. The framework must distinguish what is ideal, what is feasible, and what is required before public authority is claimed.

14.11 Objection 11: complexity defeats stable constraints

Chaotic, adaptive, and reflexive systems can change as they are studied. A constraint that is stable today may fail after an intervention. Does this destroy realism?

No. It changes the ontology of the claim. Stability can be local, dynamic, or conditional. A constraint may be stable within a regime and fail at a threshold. A social rule can persist until actors coordinate to change it. The scientific task is to model the conditions of stability and the transitions between regimes.

Revision. RCR replaces the demand for universal invariance with a search for regime-relative invariance and explicit transition conditions.

14.12 Objection 12: the framework risks relativism

If all claims are scoped and fallible, perhaps no claim is objectively true.

Fallibility is not relativism. A claim can be true even when it is revisable. The distinction is between the truth of a claim and the certainty available to an investigator. Scope is not a retreat from reality; it is a condition for honest reference. A map can be correct about roads without being a complete description of a country.

Revision. RCR should use calibrated language: established, well-supported, provisional, speculative, and rejected. These labels track warrant, not the existence of multiple opinions.

14.13 Objection 13: the framework risks scientism

A framework that gives science a theory of objective knowledge may be tempted to treat scientific methods as the standard for every question. RCR rejects that extension. Ethical obligations, meanings, political legitimacy, and aesthetic value require forms of reasoning that are not reducible to empirical science. Science can inform these questions without deciding them alone.

Revision. The framework's jurisdiction is world-directed empirical and formal inquiry. It does not claim that every rational question is a scientific question.

14.14 The revised framework

After these objections, RCR can be stated more cautiously:

Revised RCR. Scientific objectivity is the graded, fallible achievement of claims and institutions that stabilize relevant constraints through traceable measurement, discriminating comparison, intervention or causal reconstruction where possible, reflexive criticism, and scope-honest revision. The achievement warrants layered commitments to a mind-independent reality, but never an unqualified final inventory of its contents.

This statement is less dramatic than a promise of certainty. It is more useful because it identifies what can improve. A theory can become more objective by improving its measurements, alternatives, mechanisms, or institutional correction capacity even before its ontology is settled.

Chapter summary

The strongest objections do not destroy RCR; they constrain it. Constraints must be typed. Convergence must account for correlated failures. Intervention must be theory-aware. Bayesian updating requires prior sensitivity. Social criticism is a condition of access, not a substitute for world resistance. Standpoint claims require testing. Values should be transparent. The commitment ladder is a reporting device, not a fixed scale. Requirements should be proportional to stakes, and complex systems require regime-relative invariance. The revised framework is realist, fallibilist, reflexive, and non-scientistic.

Objections and discussion.

1. Which objection poses the deepest challenge to the possibility of objective science?
2. How should evidence requirements vary with the stakes of a claim?
3. Is a framework stronger or weaker when it explicitly limits its jurisdiction?

Further reading. Compare the objections with van Fraassen (1980), Stanford (2006), Fine (1984), Longino (1990), Douglas (2009), Mayo

(1996), and the literature on causal inference and scientific objectivity.

Conclusion: A Programme for Objective Inquiry

Chapter guide. The conclusion condenses the book into a systematic statement. It restates the ontology, epistemology, inferential rules, criteria of objectivity, and limits of RCR. It then lists unresolved problems, possible empirical implications, and a research programme for testing and extending the framework.

The problem of objectivity is not solved by finding a perfect foundation. Science has no access to the world without concepts, instruments, models, language, and institutions. Nor is it solved by abandoning the world. The fact that inquiry is situated does not make every description equally good. Scientific knowledge is possible because reality is not indifferent to our actions. It pushes back against some classifications, breaks some models, makes some interventions fail, and supports others.

Reflexive Constraint Realism names an attempt to understand that possibility. The theory is realist because it affirms a mind-independent, generative world. It is constraint-based because it treats objective knowledge as the stabilization of relations that resist arbitrary reinterpretation. It is reflexive because it includes the social and conceptual conditions under which evidence is made, criticized, and revised.

The systematic statement

Ontology

1. Reality exists independently of human concepts and observations.
2. Reality is generative: entities, processes, structures, and institutions produce effects and restrict possibilities.
3. Reality is layered: higher-level systems can be physically realized while retaining scale-specific causal and explanatory organization.

4. Scientific entities earn graded commitment through boundary, difference, convergence, and intervention tests.
5. Structures and laws are real to the extent that relevant relations remain stable across specified transformations and interventions.
6. Social facts are relationally real when rules, recognition, material arrangements, and sanctions give them causal force.
7. Scientific ontology is open and revisable; no current theory warrants a complete inventory of reality.

Epistemology

1. Evidence is a relation among a claim, a target, an alternative, a measurement path, and a model of error.
2. Observation is mediated but not arbitrary. It becomes objective through traceability, calibration, and independent constraint.
3. Belief should update in proportion to evidence that discriminates among alternatives. Bayesian probability is one formal tool; error analysis, intervention, mechanism, and criticism are others.
4. Causal knowledge requires stable change-relations, supported by randomization, intervention, mechanism, counterfactual analysis, or a defensible combination.
5. Models represent selectively. Their usefulness can exceed literal truth, but every model must state its idealizations and transport conditions.
6. Social criticism can strengthen objectivity when it is public, competent, independent, and capable of changing the inquiry.
7. Values may shape questions, concepts, risk thresholds, and institutions. They should not decide empirical conclusions by authority alone.
8. Objectivity is graded and fallible. A claim can be well-supported without being unrevisable.

Inferential rules

The ten rules of RCR can be compressed into a sequence of questions:

1. What exactly is the target and scope?
2. What is the measurement path?
3. Which alternatives remain live?
4. Which test would most sharply discriminate them?
5. What should remain invariant and what should change?

6. Is the claim predictive, explanatory, causal, or normative?
7. What intervention boundary is being asserted?
8. What level of ontological commitment has been earned?
9. What would lower confidence?
10. Who can inspect, criticize, reproduce, or revise the result?

These questions do not produce a universal score. They create a disciplined space in which arguments can be compared and improved.

Criteria of objectivity

A claim or inquiry system is more objective when it exhibits:

- resistance from a mind-independent target;
- convergence across routes with different potential errors;
- a traceable and validated measurement path;
- interventional or counterfactual grip where possible;
- transparent models and idealizations;
- public and competent criticism;
- calibrated uncertainty and scope honesty;
- capacity to detect and repair error.

Objectivity is therefore a relation between world, method, institution, and claim. It is not a psychological feeling of neutrality.

What RCR avoids

RCR avoids naive realism by recognizing mediation and theory-ladenness. It avoids relativism by affirming that independent constraints can make one claim better supported than another. It avoids scientism by limiting scientific authority to questions its methods can constrain. It avoids dogmatism through fallibility, prior sensitivity, alternative construction, and revision requirements. It avoids a single-method mythology by allowing different sciences to realize objectivity through different profiles of constraint.

Unresolved problems

The framework leaves important problems open.

1. **Ultimate metaphysics:** Is structure more fundamental than objects, or are both abstractions from a deeper ontology?
2. **Quantum theory:** How should intervention, causality, and measurement be interpreted when physical systems cannot be assigned

classical properties in the ordinary way?

3. **Emergence:** Which higher-level constraints are genuinely autonomous, and which can be reduced in principle but not in practice?
4. **Social kinds:** How should the reality and instability of categories such as race, gender, class, disability, and mental disorder be represented without reifying or dissolving them?
5. **Theory comparison:** Can the profile of virtues be made more precise without hiding value choices in weights?
6. **Collective rationality:** What institutional structures reliably turn disagreement into correction rather than polarization?
7. **Unconceived alternatives:** How can a research community search for theories it cannot yet formulate?
8. **AI and automated inquiry:** How should the framework evaluate models whose internal representations are not interpretable to their users?

These are not minor footnotes. They mark the boundary between a defensible framework and a final philosophy.

Possible empirical implications

A philosophical framework can have empirical implications without becoming a scientific theory of everything. RCR suggests hypotheses about inquiry systems:

1. Studies that vary measurement methods with distinct error structures should produce more reliable claims than studies that repeat one pipeline.
2. Preregistered tests and public analysis plans should reduce the gap between nominal evidence and out-of-sample performance, especially where researcher degrees of freedom are large.
3. Research communities with greater access to relevant criticism should correct errors faster, after accounting for expertise and resources.
4. Claims that report explicit intervention boundaries should transport more accurately across settings than claims stated without scope conditions.
5. Models that separate predictive fit from causal and structural claims should generate fewer unjustified policy extrapolations.
6. Including affected populations in question formation should reveal

previously neglected variables, while empirical testing determines which proposed explanations survive.

These are empirical conjectures, not established facts. They can be tested by comparative studies of laboratories, journals, research fields, measurement pipelines, and policy evaluations.

A research programme for further development

A research programme inspired by RCR would have at least six tracks.

1. **Formal epistemology.** Develop decision-theoretic and Bayesian models that represent layered truth, prior sensitivity, error correlation, and graded ontological commitment.
2. **Causal methodology.** Build tools for identifying intervention boundaries, regime-relative invariance, feedback, and transportability in complex systems.
3. **Measurement studies.** Compare how independent instrument principles, calibration chains, and construct-validity tests affect convergence in natural and social sciences.
4. **Institutional experiments.** Test whether open data, preregistration, adversarial collaboration, diverse review, and protected dissent improve error correction.
5. **Historical case analysis.** Reconstruct theory changes by tracking which constraints were retained, which ontologies were abandoned, and how measurement and intervention changed.
6. **Computational inquiry.** Apply the framework to simulations, machine-learning systems, automated laboratories, and models whose useful representations are not human-readable.

The programme should not measure success only by producing a new philosophical label. It should improve scientific practice by helping researchers say what their evidence establishes, what it does not, and what would change their minds.

Final proposition

The book's final proposition is deliberately conditional:

Objective knowledge is possible when inquiry turns the world's resistance into publicly corrigible constraints. The resulting knowledge is objective not because it is produced from nowhere, but because its claims survive relevant variation, intervention, independent criticism, and the search for error. It is realist because the constraints are not made true by agreement. It is fallible because every route to them is finite and mediated.

This proposition does not guarantee that science always succeeds. It explains why science can succeed and how its failures can become resources for better knowledge. The task of philosophy is not to sanctify science or to dissolve it into sociology. It is to clarify the conditions under which scientific inquiry deserves the authority it sometimes receives.

Closing summary

Reality under inquiry is not a mirror and not a fiction. It is a field of generative constraints approached through imperfect representations. Science earns objectivity when it makes those constraints visible, tests them through independent routes, distinguishes prediction from explanation, separates values from evidence without denying their role, and maintains institutions that can find and correct error. Reflexive Constraint Realism is offered as a framework for that achievement and as a programme for discovering where it fails.

Glossary

The glossary uses the meanings developed in this book. Many terms have wider or competing meanings in the literature; the definitions below are working definitions for RCR.

Abduction.

Inference that selects a hypothesis because it explains the evidence better than relevant alternatives. It is ampliative, not deductively certain.

Accuracy.

Closeness of a measurement or representation to the target value or structure. Accuracy differs from precision.

Ad hoc adjustment.

A modification introduced mainly to protect a claim from known counterevidence without generating independent new constraints.

Bayesian updating.

Revision of probabilities using Bayes' theorem: posterior probability is proportional to likelihood times prior probability.

Calibration.

Establishing how an instrument's response relates to a standard or independently measured quantity.

Causal effect.

A difference between outcomes under specified interventions or potential treatments, for a defined unit or population.

Causal mechanism.

Organized parts, activities, and relations whose interaction produces a phenomenon.

Claim.

A proposition with a target, content, scope, alternatives, and warrant. RCR treats a claim as incomplete if these fields remain hidden.

Construct validity.

The degree to which an operation measures the construct it is intended to represent.

Constraint.

A stable restriction or tendency on what can happen under specified

conditions. RCR distinguishes material, formal, causal, organizational, and institutional constraints.

Constructive empiricism.

Van Fraassen's view that science aims at empirical adequacy and need not require belief in the truth of unobservable theoretical claims.

Counterfactual.

A conditional claim about what would happen under a condition that may not actually occur.

Correction capacity.

The ability of an inquiry system to expose, assess, and repair error.

Data.

Local records produced by measurement, observation, or collection. Data become evidence for phenomena through analysis and validation.

Demarcation.

The problem of distinguishing science, non-science, and pseudoscience. RCR treats it as a profile of epistemic performance rather than a single binary property.

Emergence.

The occurrence of stable, explanatory, or causal patterns at a level of organization that depends on, but is not simply interchangeable with, lower-level description.

Empirical adequacy.

Correct description of observable phenomena within a specified scope, whether or not a theory's unobservable posits are accepted.

Entity realism.

The position that successful intervention with a theoretical entity can support belief in the entity's reality.

Epistemic uncertainty.

Uncertainty caused by limited knowledge of parameters, models, initial conditions, or measurement processes.

Error correlation.

The extent to which apparently independent methods share a source of failure.

Error-statistical reasoning.

An approach that emphasizes what a test would probably have detected if an error or discrepancy were present.

Explanation.

An answer to a why, how, or under-what-conditions question. Explanation may be deductive, unificatory, causal, mechanistic, or model-based.

Falsificationism.

Popper's emphasis on risky claims and attempts at refutation as central to scientific rationality.

Fallibilism.

The view that any finite belief or claim may require revision, even when well supported.

Inference to the best explanation.

Comparative reasoning that favours a hypothesis because it explains evidence better than relevant alternatives.

Incommensurability.

Difficulty of direct translation or comparison between conceptual frameworks, especially across major theory change.

Invariance.

Stability of a relevant relation under specified changes in observer, instrument, representation, intervention, or context.

Intervention boundary.

The variables changed, processes allowed to respond, and background conditions held fixed in a causal claim.

Idealization.

Deliberate simplification that removes or alters features of a target in order to investigate a selected relation.

Law-like constraint.

A stable modal relation that supports counterfactual or interventional reasoning within a defined regime.

Likelihood.

The probability of evidence conditional on a hypothesis, written $\mathbb{P}(E \mid H)$. It is not the probability that the hypothesis is true.

Measurement path.

The chain from target through interaction, instrument, calibration, transformation, error model, and interpretation.

Model.

A selective representation used to reason about a target. Models may be mathematical, computational, diagrammatic, mechanistic, statistical, or material.

Normal science.

Kuhn's term for puzzle-solving research conducted within a shared paradigm.

Objectivity.

In RCR, the graded achievement of world-sensitive, publicly corrigible,

scope-honest inquiry.

Ontological commitment.

The degree and kind of belief that evidence warrants about entities, structures, processes, or social facts.

Paradigm.

A historically situated framework of exemplars, concepts, standards, instruments, and problems that organizes research.

Pragmatic adequacy.

The degree to which a representation supports reliable inquiry, prediction, control, diagnosis, or design.

Precision.

The spread of repeated measurements. Precision does not imply accuracy.

Prior probability.

A probability assigned to a hypothesis before a specified item of evidence is incorporated.

Public corrigibility.

The institutional capacity for others to inspect, criticize, reproduce, and revise a claim.

Reflexive objectivity.

Objectivity that includes explicit awareness and critique of the social, conceptual, and material conditions of inquiry.

Reflexive Constraint Realism.

The framework developed in this book: mind-independent generative reality, layered representation, graded commitment, and public correction through constraint-sensitive inquiry.

Replication.

Repeating or extending an inquiry under a specified relation of similarity in order to test stability, error, or transport.

Research programme.

A historically extended set of theoretical commitments, auxiliary hypotheses, heuristics, and problem-solving practices.

Scope condition.

A statement of the objects, population, time, scale, interventions, or background circumstances to which a claim applies.

Scientific realism.

The view that successful scientific theories can warrant belief in both observable and some unobservable aspects of reality.

Situated knowledge.

Knowledge produced from a partial social and material location whose

limits should be made explicit.

Standpoint theory.

The view that social positions can shape what is visible and can sometimes provide epistemic advantages concerning relations of power, subject to critical testing.

Structural realism.

A family of views emphasizing that scientific knowledge or reality is fundamentally relational or structural.

Theory-ladenness.

Dependence of observation, measurement, or interpretation on prior concepts, theories, training, and instruments.

Transportability.

The validity of carrying a result from one setting to another after checking shared mechanisms and changed conditions.

Truth.

In RCR, a layered relation involving correspondence, structural accuracy, interventional adequacy, and pragmatic adequacy. These layers may diverge.

Value transparency.

Explicit reporting of value choices that shape questions, concepts, error thresholds, institutional design, or action.

World resistance.

The fact that some claims fail because independent processes do not behave as the claim requires.

How to use the glossary

The glossary is not a substitute for the arguments. In particular, the terms *constraint*, *objectivity*, and *truth* acquire their specific meaning through the relations among the axioms, inferential rules, and case analyses. A reader should return to Chapters 10 and 11 when a definition appears to conflict with a familiar use.

Editorial and Factual Review

Chapter guide. This appendix records the final review standard applied to the manuscript. It distinguishes factual verification, source verification, argument review, and visual review. It also states the remaining limits of a generated scholarly manuscript so that readers can assess it responsibly.

A publication-ready philosophy book needs more than fluent prose. It needs claims that can be traced, references that exist, arguments that do not silently change level, and pages that can be read without visual distraction. This appendix documents the review categories used for this edition.

Factual and source review

The historical claims in the manuscript were checked against primary publications where available and against respected academic reference sources for broader philosophical positions. The most consequential empirical cases use the following source classes:

The manuscript avoids invented quotations and page numbers. Where a claim is an interpretation of a historical episode, the text marks it as a philosophical reading rather than presenting it as a direct quotation from the source. Publication details should still be checked against the particular edition used for a print edition or formal scholarly citation.

Argument review

The following checks were applied to the argument:

1. Major traditions are presented in their strongest recognizable form before criticism.
2. Descriptive history, established scientific evidence, interpretation, the author's original framework, and speculation are separated.

Area	Claim type audited	Source anchor
Philosophy of science	Dates, positions, and scope of major arguments	Original books and articles; Stanford Encyclopedia entries
Physics	Brownian motion, electron, ether experiment	Einstein, Perrin, Thomson, Michelson and Morley; later historical scholarship
Biology	1953 DNA structure paper and evidential role of diffraction	Watson and Crick (1953); associated primary papers and historical records
Medicine	1948 controlled streptomycin investigation	Medical Research Council (1948), British Medical Journal
Psychology	2015 replication project design and reported interpretation	Open Science Collaboration (2015), Science
Climate	AR6 synthesis and calibrated attribution language	IPCC (2023) Synthesis Report
Economics	Policy-regime change and Lucas critique	Lucas (1976)
Social science	Audit-study design and scope	Bertrand and Mullainathan (2004)

Table 14.1: Source anchors for the factual case studies.

3. The framework has an explicit ontology, epistemology, axioms, inferential rules, criteria, objections, and limits.
4. Formal notation is defined before use. Equations are used to clarify Bayesian updating, measurement, growth, and causal inference rather than to create a false appearance of certainty.
5. Causal, predictive, explanatory, pragmatic, and normative claims are not treated as interchangeable.
6. Social and feminist approaches are included as contributions to objectivity, not as decorative additions or caricatures.
7. Historical examples are used as stress tests, not as simplistic proofs of a contemporary theory.
8. Claims are scoped where transport, intervention, or measurement is

limited.

Reference review

The bibliography was checked for the existence of the cited works, author names, publication years, titles, journals or publishers, and persistent identifiers where included. The references are deliberately conservative: when edition-specific page information was uncertain, page numbers were not invented. Readers preparing a journal article, dissertation, or commercial edition should verify the exact edition and citation style required by their publisher.

A bibliography is not evidence by itself. A cited source can be interpreted badly. The manuscript therefore uses citations as anchors for claims and states when an argument is the author's own.

Language and editorial review

The manuscript was reviewed for:

- consistent use of “science”, “scientific”, “theory”, “model”, “evidence”, “truth”, and “objectivity”;
- avoidance of unexplained specialist terms where a first-principles definition is needed;
- separation of paragraphs that make different argumentative moves;
- consistent terminology for Reflexive Constraint Realism;
- removal of unsupported superlatives and claims of finality;
- clear distinction between a method's scope and its philosophical interpretation;
- discussion questions that follow from the chapter rather than repeat its headings.

The book intentionally uses a restrained tone. It argues for a position without claiming that the position has settled every debate.

Visual and typesetting review

The LaTeX project uses a 6 by 9 inch book page, restrained colour, readable margins, consistent headers and footers, and tables with fixed column logic. The cover uses a solid background and only the title, subtitle, and author name. No decorative lines, overlapping shapes, or crowded geometry are used.

The final PDF should be rendered page by page and checked for:

1. clipped or overflowing text;
2. table collisions or unreadable cells;
3. broken equations or missing glyphs;
4. bad line breaks in headings and captions;
5. inconsistent page numbering and running headers;
6. blank pages that do not serve a chapter transition;
7. index and table-of-contents links;
8. cover spacing and contrast.

This appendix is a review record, not a guarantee that no future copy-editing decision is possible. Scholarly publication includes expert review, source checking by specialists in each case domain, and updates as evidence changes. The purpose of this edition is to provide a coherent, transparent, editable manuscript that makes those next checks possible.

Review conclusion

The manuscript meets its stated design target when it is compiled with the provided source and the rendered PDF passes the visual checks above. Its central philosophical proposal is original in formulation and argumentative organization, while its historical materials are grounded in established literature. The remaining limitations are the normal limits of a single-author work spanning many sciences: domain specialists may refine examples, later research may revise empirical assessments, and the framework itself remains open to criticism.

Bibliography

Bibliography

- Anderson, Elizabeth. 2004. "Uses of Value Judgments in Science: A General Argument, with Lessons from a Case Study of Feminist Research on Divorce." *Hypatia* 19(1): 1–24. <https://doi.org/10.1111/j.1527-2001.2004.tb01266.x>.
- Armstrong, D. M. 1983. *What Is a Law of Nature?* Cambridge: Cambridge University Press.
- Batterman, Robert W. 2002. *The Devil in the Details: Asymptotic Reasoning in Explanation, Reduction, and Emergence*. Oxford: Oxford University Press.
- Bayes, Thomas. 1763. "An Essay towards Solving a Problem in the Doctrine of Chances." *Philosophical Transactions of the Royal Society of London* 53: 370–418. <https://doi.org/10.1098/rstl.1763.0053>.
- Bechtel, William, and Adele Abrahamsen. 2005. "Explanation: A Mechanist Alternative." *Studies in History and Philosophy of Biological and Biomedical Sciences* 36(2): 421–441. <https://doi.org/10.1016/j.shpsc.2005.03.010>.
- Bertrand, Marianne, and Sendhil Mullainathan. 2004. "Are Emily and Greg More Employable than Lakisha and Jamal? A Field Experiment on Labor Market Discrimination." *American Economic Review* 94(4): 991–1013. <https://doi.org/10.1257/0002828042002561>.
- Bogen, James, and James Woodward. 1988. "Saving the Phenomena." *The Philosophical Review* 97(3): 303–352. <https://doi.org/10.2307/2185445>.
- Cartwright, Nancy. 1983. *How the Laws of Physics Lie*. Oxford: Clarendon Press.
- Cartwright, Nancy. 1999. *The Dappled World: A Study of the Boundaries of Science*. Cambridge: Cambridge University Press.
- Chakravartty, Anjan. 2007. *A Metaphysics for Scientific Realism: Knowing the Unobservable*. Cambridge: Cambridge University Press.
- Craver, Carl F. 2007. *Explaining the Brain: Mechanisms and the Mosaic Unity of Neuroscience*. Oxford: Clarendon Press.

- Douglas, Heather. 2009. *Science, Policy, and the Value-Free Ideal*. Pittsburgh: University of Pittsburgh Press.
- Duhem, Pierre. 1906. *La theorie physique: son objet et sa structure*. Paris: Chevalier et Riviere. English translation, 1954, *The Aim and Structure of Physical Theory*, translated by Philip P. Wiener. Princeton: Princeton University Press.
- Einstein, Albert. 1905. "Ueber die von der molekularkinetischen Theorie der Waerme geforderte Bewegung von in ruhenden Fluessigkeiten suspendierten Teilchen." *Annalen der Physik* 17: 549–560. English translation: "On the Movement of Small Particles Suspended in Stationary Liquids Required by the Molecular-Kinetic Theory of Heat."
- Feyerabend, Paul. 1975. *Against Method: Outline of an Anarchistic Theory of Knowledge*. London: New Left Books.
- Fine, Arthur. 1984. "The Natural Ontological Attitude." In *Scientific Realism*, edited by Jarrett Leplin, 83–107. Berkeley: University of California Press.
- Franklin, Rosalind E., and Raymond G. Gosling. 1953. "Molecular Configuration in Sodium Thymonucleate." *Nature* 171: 740–741. <https://doi.org/10.1038/171740a0>.
- Frigg, Roman, and Stephan Hartmann. 2006. "Models in Science." In *The Stanford Encyclopedia of Philosophy*, edited by Edward N. Zalta. Revised 2020. <https://plato.stanford.edu/entries/models-science/>.
- Giere, Ronald N. 1988. *Explaining Science: A Cognitive Approach*. Chicago: University of Chicago Press.
- Giere, Ronald N. 2004. "How Models Are Used to Represent Reality." *Philosophy of Science* 71(5): 742–752. <https://doi.org/10.1086/425063>.
- Goldman, Alvin I. 1999. *Knowledge in a Social World*. Oxford: Oxford University Press.
- Hacking, Ian. 1983. *Representing and Intervening: Introductory Topics in the Philosophy of Natural Science*. Cambridge: Cambridge University Press.
- Haraway, Donna. 1988. "Situated Knowledges: The Science Question in Feminism and the Privilege of Partial Perspective." *Feminist Studies* 14(3): 575–599. <https://doi.org/10.2307/3178066>.
- Harding, Sandra. 1991. *Whose Science? Whose Knowledge? Thinking from Women's Lives*. Ithaca: Cornell University Press.

- Harman, Gilbert. 1965. "The Inference to the Best Explanation." *The Philosophical Review* 74(1): 88–95. <https://doi.org/10.2307/2183532>.
- Hempel, Carl G., and Paul Oppenheim. 1948. "Studies in the Logic of Explanation." *Philosophy of Science* 15(2): 135–175. <https://doi.org/10.1086/286983>.
- Hempel, Carl G. 1965. *Aspects of Scientific Explanation and Other Essays in the Philosophy of Science*. New York: Free Press.
- Hernan, Miguel A., and James M. Robins. 2020. *Causal Inference: What If*. Boca Raton: Chapman & Hall/CRC.
- Holland, Paul W. 1986. "Statistics and Causal Inference." *Journal of the American Statistical Association* 81(396): 945–960. <https://doi.org/10.1080/01621459.1986.10478354>.
- Ioannidis, John P. A. 2005. "Why Most Published Research Findings Are False." *PLoS Medicine* 2(8): e124. <https://doi.org/10.1371/journal.pmed.0020124>.
- Kitcher, Philip. 1993. *The Advancement of Science: Science without Legend, Objectivity without Illusions*. New York: Oxford University Press.
- Kitcher, Philip. 2001. *Science, Truth, and Democracy*. New York: Oxford University Press.
- Kuhn, Thomas S. 1962. *The Structure of Scientific Revolutions*. Chicago: University of Chicago Press. 2nd ed., with postscript, 1970.
- Lakatos, Imre. 1970. "Falsification and the Methodology of Scientific Research Programmes." In *Criticism and the Growth of Knowledge*, edited by Imre Lakatos and Alan Musgrave, 91–196. Cambridge: Cambridge University Press.
- Lakatos, Imre. 1978. *The Methodology of Scientific Research Programmes: Philosophical Papers, Volume 1*, edited by John Worrall and Gregory Currie. Cambridge: Cambridge University Press.
- Laudan, Larry. 1977. *Progress and Its Problems: Toward a Theory of Scientific Growth*. Berkeley: University of California Press.
- Laudan, Larry. 1983. "The Demise of the Demarcation Problem." In *Physics, Philosophy and Psychoanalysis*, edited by Robert S. Cohen and Larry Laudan, 111–127. Dordrecht: D. Reidel.
- Longino, Helen E. 1990. *Science as Social Knowledge: Values and Objectivity in Scientific Inquiry*. Princeton: Princeton University Press.

- Lucas, Robert E., Jr. 1976. "Econometric Policy Evaluation: A Critique." In *The Phillips Curve and Labor Markets*, edited by Karl Brunner and Allan H. Meltzer, 19–46. Carnegie-Rochester Conference Series on Public Policy, vol. 1. Amsterdam: North-Holland.
- Mayo, Deborah G. 1996. *Error and the Growth of Experimental Knowledge*. Chicago: University of Chicago Press.
- Medical Research Council. 1948. "Streptomycin Treatment of Pulmonary Tuberculosis: A Medical Research Council Investigation." *British Medical Journal* 2(4582): 769–782. <https://doi.org/10.1136/bmj.2.4582.769>.
- Merton, Robert K. 1942. "The Normative Structure of Science." In *The Sociology of Science: Theoretical and Empirical Investigations*, 1973, 267–278. Chicago: University of Chicago Press.
- Michelson, Albert A., and Edward W. Morley. 1887. "On the Relative Motion of the Earth and the Luminiferous Ether." *American Journal of Science* 34(203): 333–345.
- Morgan, Mary S., and Margaret Morrison, eds. 1999. *Models as Mediators: Perspectives on Natural and Social Science*. Cambridge: Cambridge University Press.
- Nagel, Ernest. 1961. *The Structure of Science: Problems in the Logic of Scientific Explanation*. New York: Harcourt, Brace & World.
- Open Science Collaboration. 2015. "Estimating the Reproducibility of Psychological Science." *Science* 349(6251): aac4716. <https://doi.org/10.1126/science.aac4716>.
- Oppenheim, Paul, and Hilary Putnam. 1958. "Unity of Science as a Working Hypothesis." In *Minnesota Studies in the Philosophy of Science*, vol. 2, edited by Herbert Feigl, Michael Scriven, and Grover Maxwell, 3–36. Minneapolis: University of Minnesota Press.
- Peirce, Charles S. 1877–1878. "The Fixation of Belief" and "How to Make Our Ideas Clear." In *The Essential Peirce*, vol. 1, edited by Nathan Houser and Christian Kloesel, 109–123, 124–141. Bloomington: Indiana University Press, 1992.
- Pearl, Judea. 2009. *Causality: Models, Reasoning, and Inference*. 2nd ed. Cambridge: Cambridge University Press.
- Perrin, Jean. 1913. *Les Atomes*. Paris: Alcan.
- Popper, Karl. 1959. *The Logic of Scientific Discovery*. London: Hutchinson. English translation of the 1934 German edition.

- Quine, W. V. O. 1951. "Two Dogmas of Empiricism." *The Philosophical Review* 60(1): 20–43. <https://doi.org/10.2307/2266637>.
- Quine, W. V. O. 1969. "Epistemology Naturalized." In *Ontological Relativity and Other Essays*, 69–90. New York: Columbia University Press.
- Reichenbach, Hans. 1938. *Experience and Prediction: An Analysis of the Foundations and the Structure of Knowledge*. Chicago: University of Chicago Press.
- Reiss, Julian, and Jan Sprenger. 2014. "Scientific Objectivity." In *The Stanford Encyclopedia of Philosophy*, edited by Edward N. Zalta. Revised 2020. <https://plato.stanford.edu/entries/scientific-objectivity/>.
- Rubin, Donald B. 1974. "Estimating Causal Effects of Treatments in Randomized and Nonrandomized Studies." *Journal of Educational Psychology* 66(5): 688–701. <https://doi.org/10.1037/h0037350>.
- Salmon, Wesley C. 1984. *Scientific Explanation and the Causal Structure of the World*. Princeton: Princeton University Press.
- Sellars, Wilfrid. 1963. *Science, Perception and Reality*. London: Routledge & Kegan Paul.
- Stanford, Kyle P. 2006. *Exceeding Our Grasp: Science, History, and the Problem of Unconceived Alternatives*. New York: Oxford University Press.
- Thomson, J. J. 1897. "Cathode Rays." *The London, Edinburgh, and Dublin Philosophical Magazine and Journal of Science* 44(269): 293–316. <https://doi.org/10.1080/14786449708621070>.
- van Fraassen, Bas C. 1980. *The Scientific Image*. Oxford: Clarendon Press.
- Watson, James D., and Francis H. C. Crick. 1953. "A Structure for Deoxyribose Nucleic Acid." *Nature* 171: 737–738. <https://doi.org/10.1038/171737a0>.
- Weber, Max. 1949. "Objectivity in Social Science and Social Policy." In *The Methodology of the Social Sciences*, translated and edited by Edward A. Shils and Henry A. Finch, 49–112. New York: Free Press.
- Woodward, James. 2003. *Making Things Happen: A Theory of Causal Explanation*. New York: Oxford University Press.
- Worrall, John. 1989. "Structural Realism: The Best of Both Worlds?" *Dialectica* 43(1–2): 99–124. <https://doi.org/10.1111/j.1746-8361.1989.tb00933.x>.

Intergovernmental Panel on Climate Change. 2023. *Climate Change 2023: Synthesis Report*. Geneva: IPCC. <https://www.ipcc.ch/report/ar6/syr/>.

Detailed Index

- abduction, 101
- accuracy, 101
- ad hoc adjustment, 101
- Bayesian updating, 101
- calibration, 101
- causal effect, 101
- causal mechanism, 101
- claim, 101
- constraint, 57, 102
- construct validity, 101
- constructive empiricism, 102
- correction capacity, 102
- counterfactual, 102
- data, 102
- demarcation, 102
- emergence, 102
- empirical adequacy, 102
- entity realism, 102
- epistemic uncertainty, 102
- error correlation, 102
- error-statistics, 102
- explanation, 102
- fallibilism, 103
- falsificationism, 103
- idealization, 103
- incommensurability, 103
- inference to the best explanation,
20, 103
- intervention boundary, 103
- invariance, 103
- laws of nature, 103
- likelihood, 103
- logical positivism, 7
- measurement path, 103
- model, 103
- normal science, 103
- objectivity, 2, 104
- ontological commitment, 104
- paradigm, 104
- pragmatic adequacy, 104
- precision, 104
- prior probability, 104
- public corrigibility, 104
- Reflexive Constraint Realism,
104
- reflexive objectivity, 104
- replication, 104
- research programme, 104
- scientific realism, 104
- scope condition, 104
- situated knowledge, 105
- standpoint theory, 105
- structural realism, 105
- theory-ladenness, 105
- transportability, 105
- truth, 105
- value transparency, 105
- world resistance, 105